

Extending Intuitionistic Logic to a Probability Calculus (in Lean 4)

Anonymous Author(s)

Abstract

Kolmogorov’s probability fixes the event algebra to be Boolean, so every event is decidable. This is a poor fit for the semidecidable properties, termination, reachability, that verification cares about. Following the Cox–Jaynes interpretation of probability as extended logic, we define a construct for extending a logic (treated as a parameter in our approach) to a calculus of probability, and specify criteria for when this works: (1) when the resulting calculus can reconstruct the original logic in the limiting case of certain (0– or 1– probability) events, and (2) when the construction, evaluated at classical logic, agrees with classical probability theory. Instantiating the construction at intuitionistic logic gives valuations on frames, complete Heyting algebras of open sets, the algebras of verifiable propositions. In this setting, we define slack, $1 - v(A) - v(\neg A)$, the classical complement rule’s deficit and boundary mass in the measure model. Slack pinpoints exactly when Bayesian conditioning, which treats the evidence as fully decided, gives an incorrect update, disagreeing with Dempster’s rule. A separate obstruction, a five-element counterexample, is why Cox’s sum rule cannot be derived and must instead be posited. Classical probability credits the slack of a semidecidable property to its failure, an event no finite observation can confirm. The whole development is machine-checked in Lean 4 over mathlib, in roughly 5,000 lines with no sorry.

Keywords: intuitionistic logic, probability, Cox’s theorem, belief functions, locales, Lean, mathlib, formalized mathematics

1 Introduction

Kolmogorov’s probability [Kolmogorov 1933] assigns numerical values representing uncertainty to events. The events to which the probability is assigned form a Boolean σ -algebra, therefore satisfying the law of excluded middle. The axioms these assignments satisfy make up the arithmetic of chance (non-negativity, normalization, σ -additivity).

This framework commits to logical axioms that computer science cannot commit to. Many properties of programs (does this computation halt, does this process reach a given state) are only *semidecidable*: a procedure confirms the property when it holds, and no procedure, in general, confirms its failure. The natural event algebra for a semidecidable property is a *frame*, the algebra of open sets of a space, and equivalently the algebraic semantics of intuitionistic logic. We show that a sharp classical readout of a semidecidable

property, one assigning each program probability 1 if it has the property and 0 otherwise, is uncomputable: computing it uniformly is exactly as hard as deciding the halting problem. This is the pointwise form of the obstruction. Chaitin’s Ω [Chaitin 1975], independently uncomputable by a classical argument, is its aggregate form, and we interpret both within the framework we develop here. The obstruction is not specific to probability, since it restates the halting problem. What it fixes is the shape of a computable credence about halting, which cannot be sharp. Classically, the remainder $1 - P(\text{halts})$ is credited to the event that the program never halts, which no finite observation confirms. On the frame of observable events, the same remainder is instead slack. In general, reasoning about the chance that a program has some property therefore calls for an event logic in which mass sits only on events that can be observed, and that logic is not Boolean. We develop a mechanized framework that allows us to study probability calculi without committing to a fixed (classical) logic, and thereby select the probability calculus aligned with the logic of the objects it reasons about.

Applying classical probability theory to non-classical logics does not always lead to the kinds of contradictions that result in fully degenerate conclusions. In this work, we make precise what the failure modes of this probability-logic mismatch are. Classical probability is not misapplied to semidecidable or resource-sensitive events, only to the wrong carrier for them. For the halting problem, the law is closure under complement. Feed a program random bits. It halts exactly when some finite string of bits makes it halt, so halting is a countable union of the events “the input begins with this string”, hence measurable, which is why Ω above is classically well-defined. But then “never halts” is measurable too, although no finite string confirms it, and additivity credits it the whole remainder $1 - P(\text{halts})$. Updating on evidence B keeps the error: the classical posterior of the Boolean complement of A exceeds that of $\neg A$ by exactly $\mu(\partial A \cap B)/\mu(B)$ (Theorem 33), mass that a frame instead leaves as slack. For linear logic, the law is idempotence of disjunction, $a \sqcup a = a$. Grading $\otimes$ by modularity then forces $v(a \otimes a) = v(a)$ (Theorem 57), so the assignment can no longer count uses. Since frames satisfy this law too, the fix is not a change of logic. Probabilistic coherence spaces [Danos and Ehrhard 2011] change the carrier, grading independent occurrences of a resource by the product $p \cdot p = p^2$ of independent trials. Classical probability is not wrong in either case. It is asked the wrong question, on an event space whose laws answer it in advance.

Probability is, as the Cox–Jaynes tradition has it, *extended logic* [Cox 1946; Jaynes 2003]. Different logics can be extended through the same construction. First, fix an event algebra, i.e. the chosen event logic in algebraic form. This algebra consists of the underlying set of events, together with the operations interpreting the logic’s connectives. For the logics compared in this paper these are conjunction, disjunction, and negation, with the law of excluded middle holding in the classical case. Then let *valuations* be the maps from events to numbers that satisfy the axioms of monotonicity, normalization, and modularity. We specify the two sufficient and necessary conditions for when the extension construct we propose defines a probability calculus (i.e. degrees of certainty rather than an arbitrary grading of events). The first (1) is that the valuations taking only the values 0 and 1, the two degrees of complete certainty, reconstruct the underlying event logic (this is what we take “extended logic” to mean). The second (2) requires that the construction, evaluated at classical logic, returns the classical theory (this ensures that the new construction is a “generalization of classical probability”). This construction treats the logic of events as a *parameter*: the choice of event algebra determines an associated probability calculus.

Kolmogorov takes the event algebra to be Boolean, and his calculus is the value of the construction at Boolean logic. There, the law of excluded middle in the underlying logic makes it possible to show that every valuation obeys the *complement rule* $P(A^c) = 1 - P(A)$. A motivating example for varying the logic parameter is quantum uncertainty calculus, whose event algebras are the orthomodular lattices [Birkhoff and von Neumann 1936; Gleason 1957]. Here, we study the construction extending intuitionistic logic, whose event algebras are frames: complete Heyting algebras, the lattices of open sets. (1) The all-or-nothing valuations are exactly the two-valued models of the logic (Section 3). (2) Complete Boolean algebras are exactly the frames satisfying excluded middle, so the classical theory embeds as a special case.

Given a valuation $v : \Omega \rightarrow [0, \infty]$ on a frame Ω , for all a, b , by modularity, $v(a) + v(b) = v(a \sqcup b) + v(a \sqcap b)$. The pseudocomplement $\neg a$ satisfies $a \sqcap \neg a = \perp$ but not, in general, $a \sqcup \neg a = \top$, so the complement rule degrades to $v(a) + v(\neg a) \leq 1$, with a gap

$$\text{slack}(a) = 1 - v(a) - v(\neg a).$$

Classically, this is the width of the Dempster–Shafer belief/plausibility interval [Dempster 1967; Shafer 1976], quantifying ignorance. In our system it instead quantifies the mass of the part of a that is neither confirmed nor refuted, the probabilistic image of the semidecidable gap above. On the opens of a space with measure μ , the slack of an open U is $\mu(\partial U)$, the mass of its topological boundary. For verification this matters concretely: knowing how likely a property is to hold no longer tells you how likely it is to fail, and the gap

between the two, the slack, is now also required to answer that question.

How the change of logic affects statistical reasoning can be stated operation by operation. The conditional $v(a \sqcap b)/v(b)$, the product rule, and Bayes symmetry hold verbatim. The classical denominator of Bayes’ theorem, $P(B) = P(B \mid A)P(A) + P(B \mid \neg A)P(\neg A)$, degrades to a lower bound whose deficit is a conditional slack. Marginalization is exact only over partitions. The complement trick $P(A) = 1 - P(\neg A)$ degrades to the interval $[v(A), 1 - v(\neg A)]$ of width $\text{slack}(A)$. Dempster’s rule and Bayesian conditioning, classically indistinguishable, give different posteriors from the same evidence unless its slack vanishes. The value of an event A itself is read as $P(\text{int } A)$, so a nonatomic distribution queried on intervals still behaves classically: an interval’s boundary is two points, carrying no mass. The two readings differ only at boundaries carrying mass and at semidecidable hypotheses.

Cox’s original theorem required repair after Halpern’s counterexample [Cox 1946; Halpern 1999; Horn 2003]. Even with this repair, the direct translation of Cox’s theorem to intuitionistic logic is *unsatisfiable*. This is because the sum rule is irreducible: a disjunction’s value is provably not a function of its disjuncts’ values, so modularity must be posited rather than derived. Weak excluded middle, an assumption strictly between intuitionistic and classical logic, does not restore the reducibility of the sum rule: the counterexample frame itself satisfies weak excluded middle.

Contributions.

1. **Weak excluded middle** (Sections 2 and 7). Global weak excluded middle on the opens of X is exactly extremal disconnectedness of X , for compact Hausdorff X exactly the Stonean spaces.
2. **Two characterizations** (Sections 3 and 5). The complement rule holds for *every* valuation iff the frame satisfies excluded middle, the exact content of Van Horn’s classical negation axiom R3 [Horn 2003]. Dempster and Bayesian conditioning agree iff the evidence has zero slack.
3. **Structure theorems** (Sections 4 and 5). On finite frames a valuation is a classical mass function on points. In general it splits into atomic and diffuse parts, and on its Boolean core it is a belief function.
4. **Cox’s theorem, dissected** (Section 6, Appendix C). The product-rule half is independent of the logic, proved by an Aczél/Hölder generator built from scratch. The sum-rule half, a theorem classically, admits no derivation here (a five-element counterexample) and must be posited.
5. **Two grounding models** (Section 7). Any classical measure, read on the opens, is a valuation whose slack is boundary mass. A sharp readout of halting is uncomputable, and on Cantor space the value of the

halting event is Chaitin's Ω . Slack is submodular yet does not determine the valuation.

6. **The monad dichotomy** (Section 8). Product valuations are unique on $\text{Low}(\mathbb{N} \times \mathbb{N})$ and provably non-unique on $\text{Set}(\mathbb{N} \times \mathbb{N})$, where Fubini fails. The countably-presented fragment composes associatively and commutatively.
7. **A first formal treatment.** Everything above is machine-checked in Lean 4 over mathlib, to our knowledge the first mechanization of non-additive probability over intuitionistic logic.

Mechanization. The development is carried out in Lean 4 [de Moura and Ullrich 2021] over mathlib [mathlib Community 2020], in roughly 5,000 lines with no sorry, and is included in full as anonymous supplementary material.¹ Two points are specific to a proof assistant. We separate *object-level* classicality, excluded middle in the frame Ω , from the *meta-level* classicality mathlib inherits through `Classical.choice`, track which results use which, and equip the one essential non-Boolean model with a computable decidability instance. The product-rule half of Cox's theorem needs an Aczél/Hölder regratuation absent from mathlib, which we build from scratch.

2 Valuations on Frames

2.1 The Event Logic

Definition 1 (frame). A *frame* is a complete lattice Ω in which finite meets distribute over arbitrary joins,

$$a \sqcap \left(\bigsqcup_{i \in I} b_i \right) = \bigsqcup_{i \in I} (a \sqcap b_i)$$

for every $a \in \Omega$ and every family $\{b_i\}_{i \in I} \subseteq \Omega$. Equivalently, it is a complete Heyting algebra. By $x \Rightarrow y$ we denote the Heyting implication in a frame, and $\neg a := a \Rightarrow \perp$ is the pseudocomplement of $a \in \Omega$. Also, a^c denotes a complement of a , if it exists, i.e. an element such that $a \sqcap a^c = \perp$ and $a \sqcup a^c = \top$.

Frames are the algebraic semantics of intuitionistic logic and the point-free abstraction of topological spaces. The same objects with morphisms reversed form the category of *locales*, so the distinction is invisible at the level of a single algebra. We say *frame* throughout and reserve *locale* for the spatial and categorical statements (points of a locale, classifying locales) where that word is the standard one. The motivating example throughout is $\Omega = \mathcal{O}(X)$, the opens of a space, with $\sqcap, \sqcup$ as intersection and union and the pseudocomplement $\neg a$ of an open given by the *interior* of its set complement. In mathlib this is the typeclass `Order.Frame`, and the Heyting operations come with the key asymmetry already in place: $a \sqcap \neg a = \perp$ holds in any Heyting algebra

¹Claude (Anthropic) was used to assist with drafting and debugging the Lean formalization.

(`inf_compl_eq_bot`), while $a \sqcup \neg a = \top$ is available only in Boolean algebras. An element with $a \sqcup \neg a = \top$ is *complemented*: it satisfies its own instance of excluded middle.

2.2 The Valuation Structure

The central object assigns a degree of certainty to each element of the frame, with inclusion–exclusion as the only interaction axiom:

Definition 2 (valuation). A *valuation* on a frame Ω is a *monotone* function $v : \Omega \rightarrow [0, \infty]$ with $v(\perp) = 0$, $v(\top) = 1$ (*normalization*), and, for all $a, b \in \Omega$, *modularity*:

$$v(a) + v(b) = v(a \sqcup b) + v(a \sqcap b).$$

```
structure Valuation (Ω : Type*) [Order.
  Frame Ω] where
  toFun : Ω → ℝ≥0∞
  map_bot' : toFun ⊥ = 0
  map_top' : toFun ⊤ = 1
  mono' : Monotone toFun
  modular' : ∀ a b, toFun a + toFun b
    = toFun (a ⊔ b) + toFun (a ⊓ b)
```

This axiom set is not ours: Weatherston's axioms for intuitionistic probability are exactly these [Weatherston 2003], and the valuations of constructive measure theory [Coquand and Spitters 2009; Vickers 2008] are their Scott-continuous instances (Section 10).

Values live in the extended non-negative reals $[0, \infty]$, and normalization confines them to $[0, 1]^{\text{Lean}}$. The conventions this choice carries, unconditional summation and truncated subtraction, are an artifact of the mechanization and are recorded in Appendix A.

2.3 Slack

Modularity applied to the disjoint pair $(a, \neg a)$ collapses to $v(a) + v(\neg a) = v(a \sqcup \neg a)^{\text{Lean}}$. So, the join need not be $\top$:

Theorem 3 (the complement bound). *For every valuation v on a frame Ω and every $a \in \Omega$, $v(a) + v(\neg a) \leq 1$.*

```
theorem Valuation.add_compl_le_one (v :
  Valuation Ω) (a : Ω) :
  v a + v a^c ≤ 1
```

This is a *derived* Dempster–Shafer sub-additivity inequality. The deficit

$$\text{slack}(a) = 1 - (v(a) + v(\neg a)) = 1 - v(a \sqcup \neg a)$$

is the valuation of the undecided region. It decomposes canonically. The pair $(\neg \neg a, \neg a)$, double negation against negation, is always disjoint. However, it is a partition only under *weak* excluded middle ($\neg a \sqcup \neg \neg a = \top$). Applying modularity to this disjoint pair splits the slack into two logically independent gaps:

Definition 4 (regular element). An element $a \in \Omega$ is *regular* (or $\neg\neg$ -stable) if $a = \neg\neg a$.

Theorem 5 (slack decomposition^{Lean}). For every valuation v on a frame Ω and every $a \in \Omega$, $\text{slack}(a) = \text{dnGap}(a) + \text{deMorganGap}(a)$, where $\text{dnGap}(a) = v(\neg\neg a) - v(a)$ is 0 if a is $\neg\neg$ -stable (regular), and $\text{deMorganGap}(a) = 1 - v(\neg\neg a \sqcup \neg a)$ is 0 if weak excluded middle holds at a .

The single Dempster–Shafer “ignorance” number is thus resolved into a regularity defect and a De Morgan defect: the first measures failure of double-negation stability, the second corresponds to a named intermediate logic (weak excluded middle). Constructively, the contrapositive of a universal hypothesis returns only its double-negated form, so for open sets, this specification of probability does not fall victim to Hempel’s raven paradox [Hempel 1945]: evidence of the white sock indoors confirms at most $\neg\neg a$, not a itself, and dnGap is exactly the gap between confirming $\neg\neg a$ and confirming a itself.

Because dnGap and deMorganGap are themselves defined from v , the apparent three unknowns $v(a)$, $v(\neg a)$, and $v(\neg\neg a)$ are not independent: they are part of a three-constraint system, yielding the following result:

Theorem 6 (slack in closed form^{Lean}). For every valuation v on a frame Ω and every $a \in \Omega$,

$$\text{slack}(a) = (v(\neg\neg a) - v(a)) + (1 - v(\neg\neg a) - v(\neg a)),$$

where $v(a) \leq v(\neg\neg a)$.

The inequality follows from monotonicity applied to $a \leq \neg\neg a$, a fact of every Heyting algebra, and shows the first summand is exactly the difference, not the floor produced by $[0, \infty]$ ’s truncated subtraction. The split into a regularity defect and a De Morgan defect follows from monotonicity and modularity alone, with no further axiom about ignorance required. This contrasts with the usual Dempster–Shafer picture, where the mass assigned to non-singleton, ambiguous sets is an extra primitive.

Theorem 6 expresses slack pointwise in terms of v . The converse, whether slack in turn determines v , is more delicate. One constraint is automatic:

Theorem 7 (slack is submodular^{Lean}). For every valuation v on a frame Ω and every $a, b \in \Omega$,

$$\text{slack}(a \sqcup b) + \text{slack}(a \sqcap b) \leq \text{slack}(a) + \text{slack}(b).$$

The proof runs v ’s own modularity twice, once at (a, b) and once at $(\neg a, \neg b)$, and combines the two results with the De Morgan laws: the equality $\neg(a \sqcup b) = \neg a \sqcap \neg b$, which holds in every Heyting algebra, and the inequality $\neg a \sqcup \neg b \leq \neg(a \sqcap b)$, monotone under v but not in general an equality. Any function proposed as the slack of some valuation must therefore be submodular.

Remark 8 (non-recovery^{Lean}). Submodularity is necessary, not sufficient: Section 7.2 exhibits two valuations on the same frame with identical slack functions.

Remark 9 (slack under weak excluded middle^{Lean}). The two gaps, dnGap and deMorganGap , computed at a single element, are logically independent. However, if weak excluded middle is imposed on the frame for every a ($\neg\neg a \sqcup \neg a = \top$ ^{Lean}), then deMorganGap vanishes identically and $\text{slack}(a) = \text{dnGap}(a)$ for every valuation and every a . This sits strictly between plain intuitionistic logic, where both gaps can be positive, and full excluded middle, where both vanish^{Lean}. A concrete instance appears in Section 7.

2.4 What Survives without Excluded Middle

Three classical pillars survive intact, and one fails in a precisely delimited way. *Disjoint additivity* is unconditional: $a \sqcap b = \perp$ implies $v(a \sqcup b) = v(a) + v(b)$. *Conditioning* is native: for $v(b) \neq 0$ the posterior $v(a \mid b) = v(a \sqcap b) / v(b)$ is again a valuation (the proof is frame distributivity), and the product rule $v(a \mid b) \cdot v(b) = v(a \sqcap b)$ and Bayes symmetry (and therefore Bayesian updating) $v(a \mid b) v(b) = v(b \mid a) v(a)$ hold^{Lean}. *Total probability* holds over any genuine partition, meaning $a \sqcap a' = \perp$ and $a \sqcup a' = \top$: $v(b) = v(a \sqcap b) + v(a' \sqcap b)$. What fails is total probability over $\{a, \neg a\}$. That family need not join to $\top$, and the conditional masses fall short of $v(b)$ by a conditional slack^{Lean}.

The classical fragment. We characterize the elements of Ω on which the classical complement rule holds. The $\neg\neg$ -stable (regular) elements, which form a Boolean algebra (by Glivenko’s theorem), are a natural candidate, but they do not suffice. Modularity constrains the frame join $\sqcup$, whereas the Boolean join of regular elements is the regularized $\neg\neg(a \sqcup b)$, for which only the inequality $v(a) + v(b) \leq v(\neg\neg(a \sqcup b)) + v(a \sqcap b)$ survives (i.e. the 2-monotonicity of Section 5). A regular but uncomplemented element, which is an open ray $(-\infty, 0)$ in the opens of $\mathbb{R}$, under an atomic measure at 0, still carries slack. The correct classical fragment is the *complemented* elements, where $v(a) + v(a^c) = 1$ holds exactly^{Lean}. On a connected locale that fragment can be as small as $\{\perp, \top\}$. In the Boolean limit, every element is complemented, the slack vanishes identically, and classical Kolmogorov axioms are recovered^{Lean}.

Mixtures. Valuations are closed under finite convex combination^{Lean}, with every axiom inherited termwise, because modularity is linear. This innocuous fact becomes a characterization in Section 4.

3 Classicality in Theory and Metatheory

Classicality enters this work twice: in the metatheory, the logic of the proof assistant itself, and in the object theory, the frame’s own logic.

Lean's meta-theory. Lean's foundations include Classical.choice, and mathlib is built on it throughout. For example, the order relation on $\mathbb{R}$ and $[0, \infty]$ has no computable Decidable instance in general, so any library lemma that case-splits on $x \leq y$ for reals resolves that split via Classical.propDecidable. A formalization that never invokes excluded middle at the level of frame elements $a \in \Omega$ can therefore still be classical through the arithmetic it borrows from the library. We track this separately from object-level classicality throughout the paper. Theorem 20 of Section 4 rests only on this ambient baseline and contains no explicit case split. The Scott-continuity recovery result (Example 19) invokes Lean's classical tactic directly, and the explicit Cox model of Section 6 goes the other way, replacing the default with an explicit computable Decidable instance to keep that one instance genuinely non-classical.

Object-level classicality: the R3 axiom. Van Horn's repaired Cox theorem [Horn 2003] axiomatizes plausibility over classical propositions. Its negation axiom, called R3, posits that the plausibility of $\neg A$ is a fixed non-increasing function S of the plausibility of A . Consistency forces the involution $S \circ S = \text{id}$, and therefore the complement rule $p \mapsto 1 - p$. We show that the involutive negation is double-negation elimination, so R3 is where classical logic enters every known Cox-style derivation. To do this, we first isolate the exact content of R3's conclusion as a predicate on valuations, then prove equivalence to excluded middle.

That predicate is *classical negation*: a valuation v has classical negation if $v(a) + v(\neg a) = 1$ for every $a \in \Omega$.

Theorem 10 (the R3 characterization^{Lean}). *For every frame Ω , every valuation on Ω has a classical negation iff $a \sqcup \neg a = \top$ for every $a \in \Omega$.*

The forward direction is a calculation. The content is the converse: wherever excluded middle fails, a valuation with genuine slack must be *constructed*. Suppose $a \sqcup \neg a \neq \top$.

Definition 11 (ideal, prime ideal). An *ideal* $J \subseteq \Omega$ is a downward closed set which is also closed under finite joins. An ideal J is *prime* if $\top \notin J$ and $a \sqcap b \in J$ implies $a \in J$ or $b \in J$.

The prime ideal separation theorem (mathlib's DistribLattice.prime_ideal_of_disjoint_filter_ideal) yields a prime ideal J containing $a \sqcup \neg a$ but not $\top$. The indicator of J 's complement,

$$v(x) = \begin{cases} 0 & x \in J \\ 1 & x \notin J, \end{cases}$$

is a valuation. Modularity in the only nontrivial case ($x, y \notin J$) is *exactly* primeness, $x \sqcap y \in J \Rightarrow x \in J \vee y \in J$. And this valuation assigns $v(a) + v(\neg a) = 0 \neq 1$.

Sharp valuations are the points of the locale. The construction above is one half of an exact correspondence. Call v *sharp* if it is $\{0, 1\}$ -valued, a state of full certainty. Sharpness characterizes exactly the prime-ideal indicators:

Theorem 12 (sharp valuations are points^{Lean}). *For every valuation v on a frame Ω , v is sharp iff v is the complement-indicator of a prime ideal.*

The correspondence is given by two mutually inverse maps: $v \mapsto Z(v) = \{x \in \Omega : v(x) = 0\}$, which computes the zero-set of a sharp valuation, and the map taking J to its complement-indicator (the valuation of Theorem 10's construction). $Z(v)$ is prime exactly because v is sharp, and the complement-indicator of $Z(v)$ recovers v itself. In locale-theoretic terms, sharp valuations recover the *points* of the locale. A prime ideal is *completely prime* when it is also closed under arbitrary (not just finite) suprema of its members. The finitely prime points of Theorem 12 sharpen to the completely prime points exactly where Scott continuity is added, recovering points of the locale in the sense used by [Coquand and Spitters 2009; Vickers 2008].

Definition 13 (directed set, Scott continuity). A subset $D \subseteq \Omega$ is *directed* if, for any finite subset $F \subseteq D$, there exists an element $d \in D$ that is an upper bound of F (taking $F = \emptyset$ shows a directed set is nonempty). A valuation v is *Scott-continuous* if $v(\bigsqcup D) = \sup_{a \in D} v(a)$ for every directed D .

Theorem 14 (Scott-continuous sharp valuations^{Lean}). *For every valuation v on a frame Ω , v is sharp and Scott-continuous iff v is the complement-indicator of a completely prime ideal.*

The characterization and the two correspondences together locate the object-level classicality precisely: whenever excluded middle fails at some a , the failure occurs at a point of the locale.

4 Valuation Representation

Every valuation on a finite frame is a probability mass function on points. On infinite frames, this can fail, and Scott continuity marks the boundary between where it still holds and where it does not.

4.1 Finite Frames Representation

Theorem 15 (Birkhoff duality [Birkhoff 1937]). *Every finite distributive lattice is isomorphic to the lattice of lower sets of a finite poset: its sup-irreducible elements form a poset P whose lower sets reconstruct the lattice, and every finite poset arises this way from its own lattice of lower sets.*

Let P be the finite poset of sup-irreducible elements of a finite frame. Then, by Birkhoff duality (OrderIso.lowerSetSupIrred), every finite frame is the lattice of *lower sets* of P . We fix the following notation, used throughout the rest of the paper. A *lower set* $L \subseteq P$ satisfies $p \in L, q \leq p \Rightarrow q \in L$, and an *upper set* is defined dually.

Lower sets are always denoted L , and upper sets U (opens of a topological space are still U, V, W , as in traditional usage). The *principal lower set* at p is $\downarrow p = \{q : q \leq p\}$, the *strict principal lower set* at p is $\Downarrow p = \{q : q < p\}$, and the *principal upper set* at p is $\uparrow p = \{q : q \geq p\}$. We prove the finite representation theorem below on the lower sets of P . Transporting it across the isomorphism above then gives the theorem for every finite frame, not just lower sets^{Lean}. Let us define the *mass* of a point as the discrete derivative of v^{Lean} :

$$\text{mass}(p) = v(\downarrow p) - v(\Downarrow p).$$

The masses recover v and sum to 1:

Theorem 16 (finite representation^{Lean}). *For a finite P , any valuation v on $\text{Low}(P)$, and any $L \in \text{Low}(P)$, $v(L) = \sum_{p \in L} \text{mass}(p)$, and $\sum_p \text{mass}(p) = 1$.*

The proof works by induction on $|L|$: repeatedly remove a maximal point p from L , and use modularity applied to $(L \setminus \uparrow p, \downarrow p)$ to show each removal changes v by exactly $\text{mass}(p)$. The mass function is a classical probability mass function on P^{Lean} . And since lower sets are exactly the opens of the Alexandrov topology, where interior is the identity, the theorem also reads $v = \mu \circ \text{int}$, i.e. every valuation on a finite frame is the classical measure μ applied after taking the interior. In the finite case, this is the converse of the measure-theoretic bridge of Section 7. The intuitionistic mass also agrees with the classical notion of the mass of a point: reading $v = \mu \circ \text{int}$,

$$\text{mass}(p) = \mu(\downarrow p) - \mu(\Downarrow p) = \mu(\{p\}),$$

since $\downarrow p$ is the disjoint union of $\Downarrow p$ and $\{p\}$ and μ is defined on all subsets of P .

Theorem 16 represented v as a sum of point masses. The dual view represents v as a mixture of sharp *point-valuations*. For each $p \in P$, let δ_p be the valuation with $\delta_p(L) = 1$ if $p \in L$, and $\delta_p(L) = 0$ otherwise:

Theorem 17 (mixture characterization^{Lean}). *On a finite frame, every valuation v satisfies $v = \sum_p \text{mass}(p) \cdot \delta_p$.*

Together with the fact that any mixture of valuations is again a valuation^{Lean}, this theorem characterizes the valuations on a finite frame exactly: they are the finite convex mixtures of point-valuations, closed under mixing because modularity is linear. Section 6 shows modularity cannot instead be *derived* from more primitive axioms. It must be posited outright, and doing so is exactly what makes every valuation a mixture of sharp, all-or-nothing certainties.

4.2 Infinite Boundary Representation via Scott Continuity

Both halves of the finite theorem fail for infinite frames:

Example 18 (obstruction^{Lean}). On $\text{Low}(\mathbb{N})$, the indicator v of $\top$ is a valuation with $\text{mass}(p) = 0$ for every point p , and $v(\top) = 1$.

No point carries any of the mass. The failure is one of approximation: $v(\top)$ is not the supremum of $v(\downarrow n)$ over the finite stages. The $\top$ -indicator of Example 18 is not Scott-continuous^{Lean}, and Scott continuity is exactly what restores the representation:

Example 19 (recovery^{Lean}). Let v be a valuation on $\text{Low}(\mathbb{N})$. If $v(\top) = \sup_n v(\downarrow n)$ then $v(\top) = \sum_n \text{mass}(n)$.

In general, on any poset all of whose principal lower sets are finite, Scott continuity implies zero diffuse part without additional preconditions. However, subtraction in $[0, \infty]$ is truncated at 0, so if $\sum_p \text{mass}(p)$ were greater than $v(\top)$, the expression $v(\top) - \sum_p \text{mass}(p)$ would still evaluate to 0. Theorem 20 rules out $\sum_p \text{mass}(p) > v(\top)$ unconditionally, so the subtraction below is never truncated:

Theorem 20 (general decomposition^{Lean}). *For every valuation v on $\text{Low}(P)$ and an arbitrary poset P , $\sum_p \text{mass}(p) \leq v(\top)$.*

By Theorem 20 the remainder $v(\top) - \sum_p \text{mass}(p)$ is a genuine difference. Every valuation thus splits unconditionally into an *atomic* part carried by points ($\sum_p \text{mass}(p)$, which is always a sub-probability), and a *diffuse* (nonnegative) remainder $v(\top) - \sum_p \text{mass}(p)$. This is a constructive Lebesgue-style decomposition. The finite and Scott cases are the purely atomic extremes, and the $\top$ -indicator is the purely diffuse one. The proof of Theorem 20 re-runs the maximal-point induction on each finite subset of P , bounding its partial sum by $v(\top)$, then takes the supremum over all finite subsets to bound $\sum_p \text{mass}(p)$ itself. Its proof does not directly invoke an excluded-middle case split. However, it relies on the (classical) decidability of the ordering of the reals (a default in mathlib). Representing the diffuse part on non-spatial frames remains the open frontier. The obstruction there is again that the frame has *too few points*, the same one as behind the undecidability of Section 7.

5 The Dempster–Shafer Reading

Sections 2–4 independently arrive at the same structure as Dempster–Shafer (DS) theory [Dempster 1967; Shafer 1976], including sub-additive complements, a belief/plausibility interval, and mass functions. This section makes the identification exact in both statics and dynamics. In each case, it locates the precise point where excluded middle makes a difference.

5.1 Total Monotonicity via Constructive Inclusion–Exclusion

The $\neg\neg$ -stable (*regular*) elements of a frame form a Boolean algebra with meet $\sqcap$ and join $\neg\neg(a \sqcup b)$ (Glivenko). Let us call the resulting algebra the *Booleanization* of this set.

Definition 21 (belief function, plausibility function). Let B be a Boolean algebra. A DS *belief function* on B is a monotone

map $\text{Bel} : B \rightarrow [0, 1]$ with $\text{Bel}(\perp) = 0$ and $\text{Bel}(\top) = 1$ that is n -monotone for every $n \geq 2$: for all $A_1, \dots, A_n \in B$,

$$\text{Bel}\left(\bigvee_{i=1}^n A_i\right) \geq \sum_{\emptyset \neq T \subseteq \{1, \dots, n\}} (-1)^{|T|+1} \text{Bel}\left(\bigwedge_{i \in T} A_i\right).$$

Its *plausibility function* is the map $\text{Pl} : B \rightarrow [0, 1]$ given by $\text{Pl}(A) = 1 - \text{Bel}(\neg A)$.

We show v , restricted to its Booleanization, is such a belief function. The 2-monotone case is immediate from modularity^{Lean}. The full tower follows from the following result:

Theorem 22 (inclusion–exclusion^{Lean}). *For every valuation on every frame Ω , every finite family $a : \iota \rightarrow \Omega$, and every finite $s \subseteq \iota$, inclusion–exclusion holds with equality for the frame join:*

$$v\left(\bigsqcup_{i \in s} a_i\right) = \sum_{\emptyset \neq T \subseteq s} (-1)^{|T|+1} v\left(\bigsqcap_{i \in T} a_i\right).$$

We state this in Lean additively, with the odd-cardinality half of the signed sum set against the even-cardinality half, so that no $[0, \infty]$ subtraction occurs. The proof is by induction with modularity and frame distributivity.

Replacing the frame join $\bigsqcup_i a_i$ on the equality's other side with its double negation $\neg\neg(\bigsqcup_i a_i)$, the Booleanization's own join, yields the following inequality:

Theorem 23 (total monotonicity^{Lean}). *For every valuation on every frame Ω , every finite family $a : \iota \rightarrow \Omega$, and every finite $s \subseteq \iota$:*

$$\sum_{\emptyset \neq T \subseteq s} (-1)^{|T|+1} v\left(\bigsqcap_{i \in T} a_i\right) \leq v\left(\neg\neg\left(\bigsqcup_{i \in s} a_i\right)\right).$$

The Booleanization's join is given by $\neg\neg \circ \sqcup$, and its meets of regular elements are again regular. Interpreted in this way, Theorem 23 states exactly the ∞ -monotone tower: v restricted to its Booleanization is a belief function.

The proof of Theorem 23 from Theorem 22 follows from the monotonicity along $a \leq \neg\neg a$. Here, inclusion–exclusion does *not* fail constructively. It is exact for the intuitionistic disjunction. Applying the double-negation nucleus (a closure operator preserving finite meets) $\neg\neg$ to the join on the equality's right-hand side can only enlarge it. This follows from the fact that $a \leq \neg\neg a$ and that v is monotone, which turns the equality into the inequality of Theorem 23. The amount by which the right-hand side is enlarged is exactly the double-negation gap of Section 2.

Definition 21 is the classical notion, stated on an abstract Boolean algebra. In our setting each of its ingredients is already present: the Boolean algebra is the Booleanization, the belief function is v itself, restricted to the regular elements and evaluated against the Boolean join $\neg\neg \circ \sqcup$, and the plausibility function is $\text{Pl}(a) = 1 - v(\neg a)$. Additive probability does not become a belief function by weakening an axiom, since modularity for v is never relaxed. What changes is only the

join used to state the tower. The belief–plausibility interval is governed by the slack:

Theorem 24 (interval width^{Lean}). *For every valuation v on a frame Ω and every $a \in \Omega$, $\text{Pl}(a) - v(a) = \text{slack}(a)$.*

5.2 The Conditioning Characterization

DS theory carries two update rules that coincide classically. *Geometric* conditioning renormalizes belief, and under the DS dictionary it is the localic posterior $v(a \sqcap b)/v(b)$ of Section 2.

Definition 25 (Dempster's rule of conditioning^{Lean}). For a valuation v on a frame Ω , elements $a, b \in \Omega$, and evidence of nonzero plausibility, $v(b) \neq 0$, *Dempster's rule of conditioning* [Shafer 1976] states that

$$v(a \parallel b) = \frac{v(a \sqcup \neg b) - v(\neg b)}{1 - v(\neg b)}.$$

The hypothesis $v(b) \neq 0$ makes the denominator nonzero, since $v(b) + v(\neg b) \leq 1$ gives $v(b) \leq 1 - v(\neg b)$.

That is, the DS conditioning rule discards the mass incompatible with the evidence and renormalizes by its plausibility. Both the geometric posterior of Section 2 and Dempster's rule of Definition 25 are again valuations. Their modularity is the frame distributivity $(a \sqcup \neg b) \sqcap (a' \sqcup \neg b) = (a \sqcap a') \sqcup \neg b$. Their disagreement is given by the *excluded-middle gap* $\text{emGap}(a, b) = v(a) - v(a \sqcap (b \sqcup \neg b))$, i.e. the mass of a outside b 's instance of excluded middle:

Lemma 26 (the join–complement identity^{Lean}). *For every valuation v on every frame Ω and every $a, b \in \Omega$, $v(a \sqcup \neg b) = v(\neg b) + v(a \sqcap b) + \text{emGap}(a, b)$.*

Dempster's numerator therefore exceeds Bayes's by exactly $\text{emGap}(a, b)$. This gap depends on a , but whether it vanishes at every a depends only on the slack at the evidence b :

Theorem 27 (the conditioning characterization^{Lean}). *For every valuation v on every frame Ω and every $b \in \Omega$ with $v(b) \neq 0$, $(v(a \parallel b) = v(a \sqcap b)/v(b) \text{ for every } a \in \Omega) \text{ iff } \text{slack}(b) = 0$.*

The hypothesis $\text{slack}(b) = 0$ eliminates the gap $\text{emGap}(a, b)$ for every a and makes the two normalizers coincide, since $1 - v(\neg b) = v(b)$. Conversely, testing at $a := b$, geometric conditioning gives certainty 1, while Dempster gives $v(b)/(1 - v(\neg b))$, which is 1 only if the slack vanishes. Theorem 27 is the valuation-relative analogue of Theorem 10: both are iff statements about the same proposition, $b \sqcup \neg b = \top$.

We call the frame-level fact *static* and the valuation-level fact *dynamic*: the former is an equation in the frame itself, between $b \sqcup \neg b$ and $\top$, with no valuation involved, while the latter tests the same proposition through v , since $\text{slack}(b) = 0$ is exactly $v(b \sqcup \neg b) = v(\top)$. Dempster and Bayes therefore

coincide at b whenever v cannot distinguish $b \sqcup \neg b$ from $\top$, even if the two differ in the frame itself. In the measure model, Dempster and Bayes disagree precisely when the evidence's boundary carries mass.

6 Cox's Theorem

Cox's theorem [Cox 1946; Jaynes 2003] derives the probability calculus from conditions on plausibility. After Halpern's counterexample [Halpern 1999], its repaired forms [Arnborg and Sjödin 2000; Horn 2003] fix the calculus up to monotone regraduation. Every known derivation is classical, and Section 3 showed why: the negation axiom R3 is excluded middle. Dropping R3 leaves the product rule unaffected, since it is logic-free, and requires the sum rule to become an axiom rather than a theorem. Separately, the naive constructive statement, demanding a regraduation strictly monotone on all of $\mathbb{R}$, is unsatisfiable (Section 6.3).

6.1 A Cox Structure without Negation

Definition 28 (Cox model^{Lean}). A Cox model on a frame Ω is a plausibility assignment $\text{pl} : \Omega \times \Omega \rightarrow \mathbb{R}$, written $\text{pl}(a, b)$ for the plausibility of a given b , together with a conjunction functional $F : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$, such that $\text{pl}(\cdot, c)$ is monotone for every c , the boundary conditions $\text{pl}(\perp, c) = 0$ and $\text{pl}(\top, c) = 1$ hold for every c , the product rule $\text{pl}(a \sqcap b, c) = F(\text{pl}(a, b \sqcap c), \text{pl}(b, c))$ holds, F is associative, and $x \mapsto F(x, y)$ is strictly monotone for every $y > 0$.

The missing R3 axiom can be observed in the omission of any constraint on the relation between $\text{pl}(\neg a, b)$ and $\text{pl}(a, b)$. The definition restricts the monotonicity of F to $0 < y$ to avoid a contradiction with the fact that the boundary axioms force $F(x, 0) = 0$ for all x . An explicit model on the non-classical chain $[0, \infty]$ satisfies this constraint^{Lean}, as $a = \top$ is provably decidable.

6.2 The Sum Rule Is Irreducible

Classically, additivity is *derived*: R3 turns disjunction into negated conjunction via De Morgan, and the product rule does the rest. Constructively that route is closed, since $a \sqcup b \neq \neg(\neg a \sqcap \neg b)$ in general:

Example 29 (the sum rule is irreducible^{Lean}). Let Ω be the five-element frame of lower sets of the poset $o < a, o < b$ with a, b incomparable. Then there is a monotone, normalized assignment q on Ω that is additive on disjoint joins yet not modular, and no binary operation S on values satisfies $q(x \sqcup y) = S(q(x), q(y))$ for all x, y . Moreover, the pairs $(\downarrow a, \downarrow a)$ and $(\downarrow a, \downarrow b)$ have identical value-pairs $(\frac{1}{2}, \frac{1}{2})$ but joins valued $\frac{1}{2}$ and 1.

Unlike for a conjunction, a disjunction's plausibility cannot generally be expressed as its disjuncts' plausibilities via conditioning. So, a constructive Cox theorem *must* posit the sum-rule half. Modularity, which says that the calculus is

closed under mixing certainties, is what posits this, justified by Theorem 17. The failure is not repaired by strengthening the logic short of Boolean: the counterexample frame satisfies weak excluded middle^{Lean}, since every nonempty lower set contains the bottom point o , so the irreducibility persists under that intermediate assumption.

6.3 Cox Model Theorem

A regraduation is a map $g : \mathbb{R} \rightarrow [0, \infty]$ with $g(0) = 0$. No such map is strictly monotone on all of $\mathbb{R}$, since strictness would force $g(-1) < g(0) = 0$, impossible in $[0, \infty]$. The repair (strictness only on $[0, 1]$, where plausibilities live) is exactly the same domain restriction that Halpern's critique of the classical theorem turns on. The corrected statement is:

Theorem 30 (the constructive Cox theorem^{Lean}). *For every frame Ω , every Cox model on Ω whose unconditional plausibility is modular regraduates to a valuation: there exist a valuation v on Ω and a strictly monotone function g on $[0, 1]$ with $g(0) = 0$, $g(1) = 1$, and for all $a \in \Omega$, $v(a) = g(\text{pl}(a, \top))$.*

The regraduation for the sum-rule half is straightforward ($g = \text{ENNReal.ofReal}$). The product-rule half, showing F regraduates to multiplication, rests on Aczél's associativity theorem [Aczél 1966], a statement of pure analysis whose type never mentions Ω . `mathlib` has no Aczél theorem, no Hölder-style embedding for ordered semigroups on intervals, and no Cauchy functional equation under monotonicity, so we build the analytic core from scratch, obtaining a normalized, strictly monotone multiplicative generator on the positive cone^{Lean}. The construction also weakens Hölder's classical hypotheses [Hölder 1901]: the Archimedean property is *derived* from continuity and positivity rather than assumed, and the interval form of divisibility follows from the intermediate value theorem^{Lean} applied to a plain divisibility hypothesis^{Lean}. The constructive/classical divide is therefore *entirely* in the sum rule: nothing in the product-rule apparatus mentions Ω . Appendix C gives the construction. The residual analytic gap is discussed in Section 11.

7 Two Grounding Models

Two concrete models demonstrate that our abstract theory is non-degenerate. One shows every classical measure *contains* an intuitionistic probability. The other shows the intuitionistic theory cannot collapse back without solving the halting problem.

7.1 Measures through the Interior Operator

Gödel–McKinsey–Tarski interpret intuitionistic logic in modal S4: propositions are *open* sets, and the modal necessity operator $\Box$ of S4 is interpreted as the topological interior, written $\text{int}(\cdot)$ below [Gödel 1933; McKinsey and Tarski 1944]. Our valuations quantify that same translation:

Theorem 31 (measures as valuations^{Lean}). *Let μ be a probability measure on a space X , $A \subseteq X$, and define $v(A) = \mu(\text{int } A)$. Since $\text{int } U = U$ for open U , v restricts on the frame $\mathcal{O}(X)$ to $U \mapsto \mu(U)$, and this restriction is a valuation.*

Modularity is $\mu(U \cup V) + \mu(U \cap V) = \mu(U) + \mu(V)$. The slack of an open U is then $1 - \mu(U) - \mu(\text{int } U^c)$, which is the measure of the boundary ∂U . We can now prove that, in general, slack is exactly the boundary mass:

Theorem 32 (slack is boundary mass^{Lean}). *For any probability measure μ on a topological space X and any open U , $\text{slack}(U) = \mu(\partial U)$.*

The Heyting complement of U in $\mathcal{O}(X)$ is the interior of its set complement. So, the undecided region, the complement of $U \sqcup \neg U$, is exactly $\overline{U} \setminus U = \partial U$. The general sub-additivity inequality of Section 2 still holds here^{Lean}. This identification of slack with boundary mass also exhibits the modal identification $\text{Bel} = P(\Box \cdot)$ of the DS literature [Ruspini 1986] arising from the logic itself. Theorem 16 is the converse in the finite case (in the general case, the converse is the open problem, see Section 4).

The same decomposition locates a failure of classical probability calculus in conditioning, that is, in updating a belief about an open A on evidence given by an open B . The Boolean complement of A splits into its Heyting complement and its boundary, $A^c = \text{int } A^c \sqcup \partial A$, so a classical posterior that treats A as confirmed-or-refuted overstates the refutation by the boundary mass the evidence retains:

Theorem 33 (classical overconfidence in refutation^{Lean}). *Let μ be a probability measure on a topological space X , let v be its restriction to $\mathcal{O}(X)$, and let A, B be opens with $\mu(B) \neq 0$. Then*

$$\frac{\mu(A \cap B)}{\mu(B)} = v(A \mid B) \quad \text{and} \quad \frac{\mu(A^c \cap B)}{\mu(B)} = v(\neg A \mid B) + \frac{\mu(\partial A \cap B)}{\mu(B)}$$

The excess term is the conditional slack of A given B , the mass that the evidence leaves neither confirmed nor refuted. The classical posterior does not discard this mass. It credits all of it to A^c , an event no finite observation confirms. For the halting event of Section 7.3 with $B = \top$, the classical credence in non-termination is $1 - p$ while the constructive one is 0, and the difference $1 - p$ is exactly the slack.

To characterize weak excluded middle spaces, we give the (standard) definitions required:

Definition 34 (extremally disconnected). A topological space X is *extremally disconnected* if the closure of every open set is again open.

Definition 35 (Stone space, Stonean space). The *Stone space* of a Boolean algebra B is the space of ultrafilters on B , topologized by the clopen basis $\{U : b \in U\}$ for $b \in B$. *Stone duality* [Johnstone 1982] identifies Boolean algebras with exactly the compact, Hausdorff, totally disconnected spaces arising this way. A *Stonean space* is a compact Hausdorff

space that is also extremally disconnected. Equivalently, it is the Stone space of a *complete* Boolean algebra.

Remark 36 (weak excluded middle on opens^{Lean}). Section 2 noted that imposing weak excluded middle on the whole frame collapses slack to pure double-negation instability. For $\Omega = \mathcal{O}(X)$, this global condition holds exactly when X is extremally disconnected.

An immediate corollary gives the characterization:

Corollary 37 (Stonean characterization^{Lean}). *Given a compact Hausdorff space X , weak excluded middle on $\mathcal{O}(X)$ holds exactly when X is a Stonean space.*

7.2 Slack Does Not Determine the Valuation

Submodularity (Theorem 7) is a necessary condition on slack. It is not sufficient: distinct valuations, each satisfying every axiom a valuation must satisfy, can share a slack function.

Example 38 (boundaries in the three-point space^{Lean}). Let $\text{ThreePt} = \{p_A, p_B, p_Z\}$ carry the topology generated by the two singletons $\{p_A\}$ and $\{p_B\}$, so its opens are exactly $\emptyset$, $\{p_A\}$, $\{p_B\}$, $\{p_A, p_B\}$, and the whole space. For every open U , $\partial U \subseteq \{p_Z\}$.

This is a fact about the topology alone, independent of any measure: p_Z has no proper open neighbourhood, so it is the only point any open's frontier can contain.

Example 39 (non-recovery^{Lean}). There exist valuations $v \neq v'$ on $\mathcal{O}(\text{ThreePt})$ whose slack functions are identical: $\text{slack}_v = \text{slack}_{v'}$ as functions.

For weights $a, b, c \geq 0$ with $a + b + c = 1$, the point measure $\mu_{a,b,c} = a \delta_{p_A} + b \delta_{p_B} + c \delta_{p_Z}$, read through Section 7.1's bridge, gives a valuation $v_{a,b,c}$ on $\mathcal{O}(\text{ThreePt})$. By Example 38 and Theorem 32, its slack at an open U is $\mu_{a,b,c}(\partial U)$, equal to c when $p_Z \in \partial U$ and to 0 otherwise: a quantity depending only on c , not on how the remaining mass splits between a and b . So $v_{1/4,1/4,1/2}$ and $v_{1/8,3/8,1/2}$ have identical slack functions while disagreeing at $\{p_A\}$, where they take the values $1/4$ and $1/8$. Both are fully modular, monotone, and normalized: nothing in Theorem 7 or its proof distinguishes them. Weak excluded middle does not rescue recoverability. If anything, the fully classical case is the worst one for it. When X is extremally disconnected, the De Morgan gap vanishes (Section 7.1) and slack degenerates to the double-negation gap alone, but on a space where excluded middle holds outright that gap vanishes too, so slack is identically 0 for every valuation. The zero slack function is then shared by every probability measure on X , which is the most extreme failure of recovery available.

7.3 The Computability Guard

The theory developed here is non-degenerate only if there are spaces whose slack does not vanish identically. Computability supplies them:

Definition 40 (Sierpiński topology, semidecidable proposition). The *Sierpiński space* $\mathbb{S}$ is the type `Prop` of propositions carrying the topology with exactly three opens: $\emptyset$, the set $\text{halts} = \{q \mid q\}$ of true propositions, and $\mathbb{S}$ itself [Escardó 2004; Rosolini 1986]. A proposition is *semidecidable* if it is an open in this topology.

The topology encodes verifiability. A point of $\mathbb{S}$ is a proposition, say “this computation halts.” The one nontrivial observation available is membership in halts , established by producing a proof (running the computation to termination). No open captures the complementary observation, and on the frame $O(\mathbb{S}) = \{\perp, \text{halts}, \top\}$ that Σ_1 asymmetry is a theorem:

Theorem 41 (the Σ_1 asymmetry^{Lean}). $\neg \text{halts} = \perp$.

Every nonempty open of $\mathbb{S}$ contains `True`, so $\perp$ is the only open disjoint from halts . Thus, no nonempty open of $\mathbb{S}$ is evidence that the computation does not halt.

We now instantiate the measure-to-valuation construction of Theorem 31 at $X = \mathbb{S}$, allowing us to compute the slack:

Example 42 (halting slack^{Lean}). Let $\mu = p \cdot \delta_{\text{True}} + (1-p) \cdot \delta_{\text{False}}$, and let v be the resulting valuation. Then

$$\text{slack}(\text{halts}) = 1 - p = \mu(\partial\{q \mid q\}) = \mu\{\text{False}\}.$$

Since halts is already open, $v(\text{halts}) = \mu(\text{halts}) = p$. The DS “ignorance” of Section 5 is here the mass of the single point `False`, the one outcome no finite observation can ever resolve. This probability is a fact about the model, not about anyone’s ignorance.

Chaitin’s Ω . The same bridge scales to algorithmic information theory. Fix a universal prefix-free machine U , *prefix-free* in that no halting program of U is a proper prefix of another, and *universal* in that it simulates every other prefix-free machine. Give 2^ω , the space of infinite bit strings, the fair-coin measure μ . The event H of having a halting program of U as a prefix is open, since halting is established by a finite prefix (Definition 40), and prefix-freeness makes H a disjoint union of cylinders, so $v(H) = \mu(H) = \sum_{p \text{ halts}} 2^{-|p|}$: exactly Chaitin’s halting probability Ω [Chaitin 1975]. Its slack $\mu(\partial H)$ is the mass of strings approximated arbitrarily well by halting strings while never having a halting program as a prefix, and Ω ’s classical uncomputability is the aggregate form of the pointwise obstruction above. The measure-theoretic core is mechanized on Cantor space with `mathlib`’s infinite product measure: basic cylinders are clopen with measure $2^{-|w| \text{Lean}}$, and a *finite* prefix-free set of words has pairwise disjoint cylinders of total measure $\sum_w 2^{-|w| \text{Lean}}$, the Kraft-inequality content underlying Ω . The universal machine itself, and the countable extension of the finite Kraft sum, are computability theory beyond our scope.

Definition 43 (halting problem). For a code c and input n , the *halting problem* is the predicate $H(c, n) =$ “the machine with code c halts on input n ,” (`ComputablePred`).

Slack cannot be eliminated by simply declaring a classical answer for each code. The family of valuations defined below attempts to do this:

Definition 44 (sharp classical readout). For a given n , the *sharp classical readout* is the family of $\{0, 1\}$ -valued valuations $\rho_n(c) \in \text{Valuation}(O(\mathbb{S}))$, with $\rho_n(c)(\text{halts}) = 1$ if $H(c, n)$ holds, and $\rho_n(c)(\text{halts}) = 0$ otherwise.

The readout is sharp but not slack-free. Since $\neg \text{halts} = \perp$ (Theorem 41), a code that does not halt receives slack $1 - 0 = 1$, the maximum. Sharp here means that the classical answer is known, not that the complement rule holds on $O(\mathbb{S})$.

However, this family is provably not computable:

Theorem 45 (the sharp readout is uncomputable^{Lean}). For every n , the predicate $c \mapsto (\rho_n(c)(\text{halts}) = 1)$ is not computable.

The proof identifies the predicate with $H(\cdot, n)$ and applies the undecidability of the halting problem (`ComputablePred.halting_problem`).

No route back to a classical, decidable treatment survives contact with computability. Excluded middle guarantees that $H(c, n)$ is true or false for any fixed c , but that guarantee is free: it names no algorithm and exhibits no witness for which disjunct holds. Demanding an actual witness, uniformly in c , is exactly ρ_n ’s sharp readout, and Theorem 45 shows no algorithm computes it. Aggregating over the whole family of halting programs, rather than reading off one instance, does not help either: Ω is a perfectly well-defined classical probability, yet Chaitin showed it uncomputable by an independent argument turning on algorithmic randomness [Chaitin 1975]. Neither the pointwise nor the aggregate route gives classical probability a way to compute a solution to the halting problem, and each is blocked by its own uncomputability result.

We can conclude that slack is not a defect of the constructive theory. On $O(\mathbb{S})$, forcing the complement rule at halts leaves the single value $v(\text{halts}) = 1$, so every code is assigned certain termination. That assignment is computable but false of every non-halting code. Any assignment that avoids this error, giving $v(\text{halts}) < 1$ to each non-halting code, has positive slack at every such code, and the one that is moreover sharp is the uncomputable readout of Theorem 45. For the sum rule, this guard plays the role the explicit Cox-model instance plays for the product rule (Section 6): each axiom set in the paper comes with non-trivial models satisfying it.

8 Product Valuations and Kernel Composition

The structure through which classical probability supports sequential statistical reasoning is the *Giry monad*:

Definition 46 (Giry monad [Giry 1982]). The *Giry monad* G sends a measurable space X to the space $G(X)$ of probability

measures on X . Its unit sends a point $x \in X$ to the point mass at x , and its multiplication sends a measure M on $G(X)$ to the measure $A \mapsto \int m(A) dM(m)$.

Probabilistic reasoning then composes in stages, and, when the monad is commutative, independent stages compose in either order. The question is whether the intuitionistic calculus supports the same structure.

A *frame homomorphism* is a map preserving finite meets and arbitrary joins. Pushforward along a frame homomorphism $f : \Omega \rightarrow \Omega'$ sends a valuation v on Ω' to $a \mapsto v(f(a))$ on Ω , and this preserves the valuation axioms, making $\text{Val}(-)$ a functor on locales, the localic counterpart of the functor part of the Giry monad [Giry 1982]. A frame homomorphism is the opposite of a locale map, so the functor laws are checked against identity and composition of frame homomorphisms^{Lean}. The point valuations δ_p (Section 4) provide the candidate (monad) unit. Theorem 16 shows that, on finite frames, Val agrees (on objects) with the classical finite-distribution monad on points. Flattening a valuation of valuations is integration, and integration in two variables requires product valuations. A *kernel* is a map assigning a valuation to each point of a locale, giving a conditional distribution. By Kock's identification of commutative monads with Fubini theorems [Kock 2012], kernels compose independently of order exactly when Val is a commutative monad.

A valuation u on a product frame is a *product* of valuations v and w when $u(A \times B) = v(A) \cdot w(B)$ for all rectangles $A \times B$. Let $\text{Low}(\mathbb{N} \times \mathbb{N})$ be the frame of lower sets of the grid $\mathbb{N} \times \mathbb{N}$, ordered component-wise. On this frame, products are unique:

Theorem 47 (product uniqueness^{Lean}). *Let v, w be valuations on $\text{Low}(\mathbb{N})$. Any two valuations on $\text{Low}(\mathbb{N} \times \mathbb{N})$ that are products of v and w are equal.*

By the inclusion–exclusion equality, two valuations agreeing on a $\sqcap$ -closed family agree on all finite joins of its members^{Lean}. The rectangles are closed under $\sqcap$, and in $\text{Low}(\mathbb{N} \times \mathbb{N})$ every lower set is a *finite* union of rectangles^{Lean}, so agreement on rectangles is agreement everywhere. Valuations carry no continuity axiom (Section 4), so the σ -additivity driving the classical uniqueness argument for product measures has no analogue here. The order structure of the frame does that work instead.

Theorem 48 (product non-uniqueness^{Lean}). *On the powerset frame $\text{Set}(\mathbb{N} \times \mathbb{N})$ there exist valuations v, w and two distinct products of v and w .*

An *ultrafilter* on $\mathbb{N}$ is a collection of subsets, closed under intersection and superset, containing exactly one of S and its complement for every $S \subseteq \mathbb{N}$. It is *free* when it contains no finite set. The map assigning 1 to the members of an ultrafilter and 0 to every other subset is a valuation^{Lean}. A rectangle $S \times T$ belongs to an ultrafilter on the product exactly when S and T belong to its images under the two coordinate projections^{Lean}. A free ultrafilter $\mathcal{U}$ on $\mathbb{N}$ yields two ultrafilters

on $\mathbb{N} \times \mathbb{N}$. A set belongs to the former when, for $\mathcal{U}$ -most first coordinates, its slice contains $\mathcal{U}$ -most second coordinates. A set belongs to the latter under the same test with the coordinates in the opposite order.

The two ultrafilters' valuations agree on *every* rectangle yet differ on the upper triangle $\{(m, m') \mid m < m'\}$ ^{Lean}. Integrating the same marginal data in the two orders thus disagrees. Fubini fails for valuations, as it does for non-additive integration [Denneberg 1994; Ghirardato 1997]. By Kock's correspondence, no canonical commutative multiplication exists for unrestricted valuations. This half of the dichotomy is not specific to intuitionism, since powerset-frame valuations are exactly the classical finitely additive probability measures [Rao and Rao 1983]. Dropping continuity is what admits the non-uniqueness. A monad therefore requires restricting either the frames or the valuations.

A first step in the second direction is to restrict Val to the valuations given as countable mixtures. For these, we can demonstrate that the kernels compose independently of order, the property that Theorem 48 refutes for unrestricted valuations:

Proposition 49 (countable kernel composition). 1. A countable mixture of valuations is a valuation^{Lean}.
2. A mixture of mixtures flattens to a mixture with product weights, which is discrete Chapman–Kolmogorov^{Lean}.
3. The one-point mixture is the identity^{Lean}.
4. Iterated countable mixing is order-independent^{Lean}.

Restricted to countable-mixture presentations, the data of a monad are all in place: the category of locales, Val restricted to countable mixtures as the functor, the point valuations as the unit, and countable mixing as the multiplication. Proposition 49 verifies exactly the laws this needs: item 1 keeps the multiplication inside the fragment, item 2 is its associativity, item 3 its unit law, and item 4 its commutativity. What is not yet settled is whether this descends to valuations themselves, since a valuation can admit more than one countable-mixture presentation, and the diffuse part of Section 4 is not captured by any countable presentation at all.

9 Valuations Beyond Frames

Definition 1 fixes the event algebra to be a frame: a complete lattice in which finite meets distribute over arbitrary joins $\bigsqcup_{i \in I} a_i$, the operation Heyting implication $\Rightarrow$ (the right adjoint to $a \sqcap (-)$) is built from. Definition 2's four axioms, normalization at $\perp$ and $\top$, monotonicity, and modularity $v(a) + v(b) = v(a \sqcup b) + v(a \sqcap b)$, use only the *binary* join $\sqcup$; none of the four mentions the *arbitrary* join $\bigsqcup_{i \in I}$. Completeness is used only where later sections evaluate a valuation on an infinite join, as in the Scott-continuity results of Section 4, or build $\Rightarrow$ and $\neg$ from it, as in Section 3's slack. Section 11 leaves open how much of the construction is tied to frames specifically, and whether some ambient category sends an event algebra to a probability calculus in general.

This section settles both questions for one fragment of the construction, the sharp/point correspondence of Section 3: it isolates the exact algebraic hypothesis that correspondence needs and shows the hypothesis cannot be dropped. It does not recover $\Rightarrow$ or $\neg$ without completeness: the adjoint defining $\Rightarrow$ is built from an arbitrary join of witnesses, so it is unavailable once completeness goes, even if only finite joins are missing. The valuations studied below are negation-free by construction, not frame valuations with completeness quietly removed.

The functor without completeness. A bounded lattice is a lattice with a least element $\perp$ and a greatest element $\top$; a bounded lattice homomorphism preserves $\perp$, $\top$, $\sqcup$, and $\sqcap$. Definition 2's axioms make sense over any bounded lattice Ω , complete or not:

Definition 50 (lattice valuation^{Lean}). A lattice valuation on a bounded lattice Ω is a monotone function $v : \Omega \rightarrow [0, \infty]$ with $v(\perp) = 0$, $v(\top) = 1$, and, for all $a, b \in \Omega$, $v(a) + v(b) = v(a \sqcup b) + v(a \sqcap b)$.

```

structure LValuation ( $\Omega : \text{Type}^*$ ) [Lattice
   $\Omega$ ] [BoundedOrder  $\Omega$ ] where
  toFun :  $\Omega \rightarrow \mathbb{R}_{\geq 0\infty}$ 
  map_bot' : toFun  $\perp = 0$ 
  map_top' : toFun  $\top = 1$ 
  mono' : Monotone toFun
  modular' :  $\forall a b, \text{toFun } a + \text{toFun } b$ 
    = toFun ( $a \sqcup b$ ) + toFun ( $a \sqcap b$ )

```

A lattice valuation on a frame is exactly a valuation in the sense of Definition 2; the two notions differ only in how much structure the ambient Ω carries.

Section 8 pulls a valuation back along a frame homomorphism and checks the assignment is functorial (identity, composition). That argument pattern-matches on $\sqcup$, $\sqcap$, $\perp$, $\top$ alone and never touches $\sqcup$:

Proposition 51 (pullback without completeness^{Lean}). For bounded lattices Γ, Ω , a bounded lattice homomorphism $f : \Gamma \rightarrow \Omega$, and a lattice valuation v on Ω , $v \circ f$ is a lattice valuation on Γ . This assignment sends the identity homomorphism to the identity and composition of homomorphisms to composition of the induced maps.

$\text{Val}(-)$ is therefore a functor on bounded lattices with bounded lattice homomorphisms, of which the frame-level functor of Section 8 is the restriction to the subcategory of frames and frame homomorphisms. Completeness contributes nothing to functoriality.

Separation needs only distributivity. Theorem 12 characterizes what a sharp valuation is, a complement-indicator of a prime ideal, but not whether there are enough of them to distinguish elements of Ω . That is a separate fact:

Theorem 52 (sharp valuations separate points^{Lean}). For every frame Ω and $a, b \in \Omega$ with $a \not\leq b$, there is a sharp valuation v with $v(a) = 1$ and $v(b) = 0$.

The proof runs the same prime-ideal separation theorem that built the counterexample of Theorem 10 (`DistribLattice.prime_ideal_of_disjoint_filter_ideal`), on the principal filter of a and the principal ideal of b rather than on $a \sqcup \neg a$ and $\top$. As a corollary, any two distinct $a, b \in \Omega$ are distinguished by some sharp valuation: Ω need not be linear, so $a \neq b$ alone does not hand back $a \leq b$ or $b \leq a$, but it forces at least one failure of $\leq$, and Theorem 52 applies in whichever direction fails. This is the cash-out of the first qualifying condition of Section 11: the certain valuations do not merely characterize a single point of the locale in isolation, together they reconstruct the order on Ω . Reconstructing the order is not the same as reconstructing the full Heyting structure: $\neg$ and $\Rightarrow$, as just noted, are adjoints built from an arbitrary join and carry no meaning over a plain distributive lattice. What this theorem recovers is the order-theoretic core of the first qualifying condition, not the negation fragment of intuitionistic logic, which stays frame-specific.

Tracing the proof of Theorem 52 shows it never invokes $\sqcup$; the only place frame structure enters is `DistribLattice.prime_ideal_of_disjoint_filter_ideal` itself, which needs distributivity and nothing beyond it:

Theorem 53 (separation over distributive lattices^{Lean}). For every distributive bounded lattice Ω and $a, b \in \Omega$ with $a \not\leq b$, there is a sharp lattice valuation v on Ω with $v(a) = 1$ and $v(b) = 0$.

Theorem 53 reproves Theorem 52 with the hypothesis `Order.Frame` weakened to `DistribLattice`, and nothing else about the argument changes. The condition doing the work in the sharp/point correspondence is distributivity; completeness is not load-bearing for either half of it.

Distributivity is not a proof artifact. The natural question is whether Theorem 53's hypothesis is a genuine dividing line or an artifact of routing the argument through a mathlib lemma named for distributive lattices. It is genuine, and the failure it guards against is total, not merely a failure of separation:

Definition 54 (the diamond^{Lean}). M_3 is the five-element lattice with a bottom z , a top o , and three pairwise-incomparable atoms a, b, c , with $x \sqcup y = o$ and $x \sqcap y = z$ for every two distinct atoms x, y .

M_3 is the smallest lattice that is not distributive: $a \sqcap (b \sqcup c) = a \sqcap o = a$, while $(a \sqcap b) \sqcup (a \sqcap c) = z \sqcup z = z$. It is modular, so this failure is specific to distributivity rather than some coarser defect of the order.

Theorem 55 (no sharp valuation on the diamond^{Lean}). No lattice valuation on M_3 is sharp.

Modularity applied to each of the three pairs of distinct atoms forces $v(x) + v(y) = 1$, since $x \sqcup y = \top$ and $x \sqcap y = \perp$. If v were sharp, $v(a), v(b), v(c) \in \{0, 1\}$ would have to satisfy all three pairwise sum-to-one equations at once, and a two-element codomain cannot support that: fixing any one value propagates through the other two equations to a contradiction with the third. The collapse on M_3 is therefore not a failure of separation, points that happen to coincide, but a failure of existence: there is no all-or-nothing valuation on M_3 to separate anything with, and the first qualifying condition of Section 11 is unsatisfiable outright.

Where this leaves the boundary. Frames are complete distributive lattices. Theorem 53 shows completeness is not needed for the sharp/point correspondence's separation half, and Theorem 55 shows distributivity cannot be given up in exchange. The event algebras on which sharp valuations reconstruct the order, the order-theoretic core of the first qualifying condition of Section 11, are exactly the distributive bounded lattices, no more and no less. The second qualifying condition and the negation-dependent machinery of Sections 3–7, slack, the complement bound, the De Morgan gap, are untouched by this section and remain tied to frames: nothing here shows, or claims, that $\neg$ or $\Rightarrow$ survives dropping completeness. This bears on the orthomodular event lattices of quantum probability discussed in Section 10: giving up distributivity, as those lattices do, does not merely change which valuations exist, it can remove sharp valuations altogether, so the construction of this paper does not, without further modification, have anything to say about that setting even at the order-theoretic level examined here. Whether some weaker notion of certainty survives on non-distributive lattices, in place of the $\{0, 1\}$ -valued sharp valuations examined here, is not addressed by this section.

A second boundary: grading a non-idempotent tensor. Section 10 notes that linear logic's quantale tensor $\otimes$ is not the lattice meet $\sqcap$ this section's valuations are built from, so Definition 50's modularity axiom does not even parse for $\otimes$ without first deciding what to write in $\sqcap$'s place. Granting the substitution provisionally, to see what it costs, shows the obstruction runs deeper than a notational mismatch:

Definition 56 (tensor valuation^{Lean}). Fix a join-semilattice Ω and a binary operation $\otimes$ on Ω . A *tensor valuation* for $\otimes$ is a function $v : \Omega \rightarrow [0, \infty]$ such that, for all $a, b \in \Omega$, $v(a) + v(b) = v(a \sqcup b) + v(a \otimes b)$.

Theorem 57 (the transplant collapses^{Lean}). For every join-semilattice Ω , binary operation $\otimes$ on Ω , and tensor valuation v for $\otimes$, if $v(a) < \infty$ then $v(a \otimes a) = v(a)$.

Setting $b := a$ in the defining identity and using only $a \sqcup a = a$, join-idempotence, already gives $v(a) + v(a) = v(a) + v(a \otimes a)$, hence the claim by cancellation. No property of $\otimes$ is used: not associativity, not a unit, not distributivity over $\sqcup$.

This is a sharper failure than Theorem 55's: M_3 shows meet-based modularity can fail to admit a sharp valuation on one uncooperative lattice, while Theorem 57 shows the $\otimes$ -based transplant admits no valuation separating a from $a \otimes a$ on any Ω and for any $\otimes$, idempotent or not. A non-idempotent $\otimes$, linear logic's whole point, is exactly what such a valuation could never distinguish from $\sqcap$'s behavior.

Example 58 (a concrete witness^{Lean}). Take $\Omega = \mathbb{N}$ with $\sqcup = \max$ and $\otimes = +$: addition is associative and commutative but not idempotent ($1 + 1 = 2 \neq 1$), so the hypothesis of Theorem 57 is not vacuous. Substituting $a = 0$, where the join's unit and the tensor's unit coincide, forces $v(n) = v(0)$ for every n : every tensor valuation for $+$ is constant.

This is not a toy chosen merely by analogy with linear logic. Every quantale, the event algebra Section 10 attaches to linear logic, is by definition a complete lattice equipped with an associative $\otimes$ that need not be idempotent, exactly the data $(\Omega, \sqcup, \otimes)$ Theorem 57 quantifies over, with the quantale's own join playing the role of $\sqcup$. Example 58 is a quantale itself ($+$ distributes over $\max$), kept only because it is small enough to check by hand: it fails idempotence for the same reason $\otimes$ does, no algebraic law identifies two copies of a resource with one.

Corollary 59 (grading and discrimination are exclusive^{Lean}). For every join-semilattice Ω , binary operation $\otimes$ on Ω , and $a \in \Omega$ with $v(a) < \infty$: no v is both a tensor valuation for $\otimes$ and satisfies $v(a \otimes a) \neq v(a)$.

This restates Theorem 57 as the dichotomy it really is. A candidate probabilistic semantics for a non-idempotent $\otimes$ must choose between being graded by Definition 50's compositional law and discriminating the one thing a non-idempotent $\otimes$ exists to discriminate, a from $a \otimes a$; it cannot keep both. The probabilistic semantics that does exist for linear logic, probabilistic coherence spaces [Danos and Ehrhard 2011], keeps discrimination: two independent copies of a resource are not scored as interchangeable with one. Corollary 59 shows the price of that choice is forced, not chosen: probabilistic coherence spaces cannot also be graded by this paper's modularity law, on pain of losing exactly the discrimination they exist to keep, so they instead evaluate proofs with a classical, Kolmogorov-style structure over $\otimes$'s type-level behavior, never a valuation of propositions in Definition 50's sense. Nothing here is a missed opportunity on either side; the dichotomy forecloses the combination before any further axiom of Definition 50 is even reached.

10 Related Work

Informal precedents. Pairing probability with intuitionistic logic is not a new idea, though the literature on it is scattered. Weatherston develops intuitionistic probability

axiomatically and defends it with Dutch-book arguments [Weatherson 2003]. His axioms match our valuation axioms exactly, with negation derived rather than assumed, so the object itself and the bound $v(a) + v(\neg a) \leq 1$ are his. What is new here is the decomposition of that bound's slack into dnGap and deMorganGap and its determinacy from v alone (Theorem 6). Paris gives an informal theorem that any modular assignment of non-classical truth values is exactly a Dempster–Shafer belief function [Paris 1994]. We mechanize the analogous fact for our intuitionistic setting: the 2-monotone case follows directly from modularity^{Lean}, and the full tower is Theorem 23. The identification $\text{Bel} = P(\Box \cdot)$ connecting belief functions [Dempster 1967; Shafer 1976] to modal logic is established [Ruspini 1986]. Section 7 shows this identification is not a separate axiom for us but a special case of our frame-level slack, recovered automatically once the frame is a topology and $\neg$ is read as the interior of the complement^{Lean}.

On the localic side, valuations on frames are standard objects of constructive measure theory, with Vickers's work and the Coquand–Spitters development as canonical references [Coquand and Spitters 2009; Vickers 2008]. These normally build in Scott continuity and target additive, Riesz-style goals. Section 4 shows Scott continuity to be a substantive assumption (it decides point-representability on the frames we examine), so our base structure omits it. Cox, Jaynes, Halpern's counterexample, and the Van Horn and Arnborg–Sjödín repairs [Arnborg and Sjödín 2000; Cox 1946; Halpern 1999; Horn 2003; Jaynes 2003] are classical throughout, and Cox's and Van Horn's derivations both lean on double-negation elimination. Section 6 proves the sum-rule half without that move, at the price of positing modularity as an axiom rather than deriving it (Example 29). Van Horn observes that dropping axiom R3 yields an interval-valued theory, exactly the belief/plausibility pair our slack separates. To our knowledge, no prior work, informal or formal, states either characterization theorem: complement rule $\iff$ excluded middle (Theorem 10), or Dempster = Bayes $\iff$ zero slack at the evidence (Theorem 27).

Prior formalizations. Two bodies of mechanized work are adjacent, and neither overlaps. ALEA [Audebaud and Paulin-Mohring 2009] formalizes a discrete distribution monad in Coq for randomized programs: additive probability, no event logic in our sense. Closest is the HoTT/Coq line of Faissole and Spitters on *valuations* over lower reals [Faissole and Spitters 2017], done constructively with Scott continuity and aimed at integration (Riesz) and a Giry-style valuation *monad*. Their valuations differ from ours in carrying Scott continuity, and their monad is the Scott-continuous instance of the question examined in Section 8. Our development drops the continuity their base structure builds in, and the two mechanizations do not otherwise overlap. On the domain-theoretic side, commutativity of the full valuations

monad on the category of directed-complete partial orders is open. A hierarchy of commutative submonads, generated from simple valuations (finite mixtures of point valuations), has been constructed instead [Jia et al. 2021]. Our countably-presented fragment, the valuations given as countable mixtures (Proposition 49), is the frame-side analogue.

Non-Boolean probability elsewhere. Quantum probability on orthomodular lattices, with Gleason's theorem [Birkhoff and von Neumann 1936; Gleason 1957], is the long-standing precedent for re-basing probability on a non-Boolean event logic. Orthomodular lattices give up distributivity itself, while our frames stay distributive by definition (Definition 1), locating all the non-classical behavior in the failure of excluded middle alone.

Linear logic marks a different kind of departure. Its event algebras are quantales, complete lattices equipped with an associative tensor $\otimes$ that need not be idempotent [Girard 1987]. The lattice reduct of a quantale is already an instance of Section 9's generalization, but $\otimes$ itself is not $\sqcap$, and a valuation grading it cannot be phrased in Definition 50's vocabulary without first deciding what modularity even means for a non-idempotent operation. The probabilistic semantics that does exist for linear logic, Girard's probabilistic coherence spaces as developed by Danos and Ehrhard [Danos and Ehrhard 2011], grades the multiplicative structure directly rather than valuing a lattice, and is not a special case of the construction studied here.

Free probability, fuzzy measures, and imprecise probability [Walley 1991] vary other parameters (independence, additivity, single-valuedness). Where Walley's coherent lower previsions posit the belief/plausibility gap directly as a coherence primitive, our slack (Section 2) is derived, forced by modularity and negation rather than assumed. Fritz's synthetic Markov-category framework [Fritz 2020] varies the ambient structure categorically rather than order-theoretically, axiomatizing conditional independence and sufficient statistics abstractly across discrete, measure-theoretic, and Gaussian probability alike. It remains unmechanized and is orthogonal to our variation in the underlying logic. Wolpert and Kinney model mathematics itself as a stochastic process, replacing binary theoremhood with a probability distribution [Wolpert and Kinney 2024]. Theoremhood is a semidecidable event, so the computability guard of Section 7 bears directly on that program: sharp credences about such events are uncomputable. Our variation is specifically the *logic*, holding the Cox-style methodology fixed. Within DS theory, the two conditionings and their inequivalence are classical [Fagin and Halpern 1991; Shafer 1976]. What is new in Theorem 27 is the exact logical characterization of their agreement, and its mechanization.

11 Conclusion

In this work, we take seriously the Cox–Jaynes reading of probability as *extended logic*, and make the extension a construction with the logic of events as a parameter. The construction fixes an event algebra (the chosen logic in algebraic form) and takes as its assignments of degrees of certainty the valuations: the maps from events to numbers satisfying monotonicity, normalization, and modularity. We specify the two conditions that qualify the result as a probability calculus: when events are certain (0 or 1 certainty), the valuations must reconstruct the underlying event logic, and the value of the construction at classical logic must be the classical theory.

The motivation for developing this extension comes from the study of program behaviour in terms of outcome probabilities. Many properties of programs are only semidecidable. However, the Boolean event algebra underlying Kolmogorov’s axioms imposes the law of excluded middle (and therefore decidability) on the structure of logical connectives applied to events. This makes it not a good fit for reasoning about events that are not decidable, such as whether an arbitrary program halts: insisting on a sharp classical readout for such an event does not decide it, but instead produces an object we show to be uncomputable, and a non-sharp classical credence assigns its remainder to non-termination, an event no finite observation confirms. The extension process we propose, when fixing its logic parameter to be intuitionistic, does not conflict with the underlying event logic of computation by construction, and is capable of reasoning about the possibility of non-termination.

Substituting the Heyting event algebra (and therefore intuitionistic logic) for the Boolean algebra representing the classical logic case changes only the consequences that excluded middle supplied. Both qualifying conditions hold at this parameter value (Sections 2 and 3), and classicality enters only through the complement rule (Theorem 10) and its dynamic counterpart, the agreement of Dempster and Bayesian conditioning (Theorem 27). The complement rule degrades to the inequality $v(a) + v(\neg a) \leq 1$. The deficit in this inequality, for which we re-purpose the term slack, is a new quantity: the value of the undecided region, given by the boundary mass in the measure model. Slack is submodular yet does not determine the valuation. The weaker logic also adds the two-gap decomposition of slack, the atomic/diffuse divide, and a belief-function reading through $\neg\neg$.

Everything else holds verbatim on frames: conditioning, the product rule, Bayes symmetry, disjoint additivity, total probability over genuine partitions, and inclusion–exclusion itself. On finite frames, every valuation is a classical mass function on the frame’s points. The product-rule half of Cox’s theorem is independent of the logic while the sum-rule half

must be posited, so the direct transcription of the classical Cox statement is unsatisfiable. Every result is machine-checked in Lean 4 over mathlib.

A recurring shape. Three of this paper’s negative results share a form worth naming: two individually reasonable demands on a valuation, jointly unsatisfiable, with the proof identifying which demand is forced to give rather than merely asserting that some construction fails. On M_3 (Theorem 55), sharpness and modularity together are unsatisfiable: a 2-valued readout cannot satisfy all three pairwise sum-to-one equations at once, and distributivity is the hypothesis that has to give. Grading a non-idempotent tensor (Corollary 59) is the same shape one level up, at the axiom itself: being graded by modularity’s compositional law and discriminating a from $a \otimes a$ cannot both hold, for any Ω and any $\otimes$, and probabilistic coherence spaces [Danos and Ehrhard 2011] are the witness that keeps discrimination and gives up gradability. Demanding a sharp, uniform, *computable* readout of a semidecidable event (Theorem 45) is the same trade-off again, one level further out: bivalence and computability cannot both hold, and the classical route keeps bivalence by fiat and gives up an algorithm. None of the three results says only that a construction fails; each names the exact pair of demands responsible and shows the failure is forced, not an artifact of how the construction happened to be set up.

Toward a monad. Section 8 lays the groundwork for a categorical exploration of the construction: Val is a mechanized functor on locales, and its countably-presented fragment composes as kernels, though whether that composition is canonical on valuations rather than their presentations remains open. A commutative monad structure, where it exists, would be the intuitionistic counterpart of the Giry monad and a step toward compositional semantics for probabilistic programs with semidecidable events. One candidate formulation exists: the space of Scott-continuous valuations is the classifying locale of a geometric theory [Coquand and Spitters 2009; Vickers 2008]. Deleting its continuity axiom still leaves a geometric theory, whose points, read classically, are exactly our valuations. Its multiplication, however, must integrate against a merely modular valuation, and the diffuse part of Example 18 is exactly what such integration fails to capture. We offer the formulation as a target, not a result. Beyond intuitionistic logic, no ambient category sending a logic to its probability calculus functorially is evident from our construction, and this question remains open.

Other open problems and future work. A natural extension of our mechanization effort involves completing the Aczél development, and mechanizing the Dutch-book theorem [Weatherson 2003]. Both require substantial additional effort and were deemed out of scope for this work. To extend our mechanization effort with preliminary monad results, we would quotient the countable presentations (Proposition 49)

into a Kleisli category of countably-presented kernels, and mechanize the classifying-locale formulation (which needs classifying locales of geometric theories, infrastructure mathlib currently lacks).

Some related mathematical problems remain open as well. The first is developing a representation of the diffuse part on non-spatial frames. The second is deriving modularity from a criterion more primitive than the mixture characterization (Theorem 17), knowing it cannot come from disjunction data alone (Example 29).

The construction method applies beyond intuitionistic logic: state the classical desiderata, delete or substitute the axioms the target logic no longer licenses, and check what survives, what splits, and what becomes unsatisfiable. Linear logic was the first case tried, and Section 9 closes it in the negative rather than leaving it open: substituting the quantale tensor $\otimes$ for $\sqcap$ in modularity is not merely difficult but unsatisfiable (Corollary 59), so no valuation in this paper's sense can grade $\otimes$'s consumption of evidence without losing the very discrimination $\otimes$ exists to carry [Danos and Ehrhard 2011]. What remains open is the next question this raises, not this instance of it: whether some compositional law other than modularity transplanted verbatim could play our construction's role for logics whose primitive connectives are not idempotent.

References

János Aczél. 1966. *Lectures on Functional Equations and Their Applications*. Academic Press, New York, NY, USA.

Stefan Arnborg and Gunnar Sjödin. 2000. Bayes Rules in Finite Models. In *European Conference on Artificial Intelligence (ECAI)*. IOS Press, Amsterdam, Netherlands, 571–575.

Philippe Audebaud and Christine Paulin-Mohring. 2009. Proofs of Randomized Algorithms in Coq. *Science of Computer Programming* 74, 8 (2009), 568–589.

Garrett Birkhoff. 1937. Rings of Sets. *Duke Mathematical Journal* 3, 3 (1937), 443–454.

Garrett Birkhoff and John von Neumann. 1936. The Logic of Quantum Mechanics. *Annals of Mathematics* 37, 4 (1936), 823–843.

Gregory J. Chaitin. 1975. A Theory of Program Size Formally Identical to Information Theory. *J. ACM* 22, 3 (1975), 329–340.

Thierry Coquand and Bas Spitters. 2009. Integrals and Valuations. *Journal of Logic and Analysis* 1, 3 (2009), 1–22.

Richard T. Cox. 1946. Probability, Frequency and Reasonable Expectation. *American Journal of Physics* 14, 1 (1946), 1–13.

Vincent Danos and Thomas Ehrhard. 2011. Probabilistic Coherence Spaces as a Model of Higher-Order Probabilistic Computation. *Information and Computation* 209, 6 (2011), 966–991.

Leonardo de Moura and Sebastian Ullrich. 2021. The Lean 4 Theorem Prover and Programming Language. In *Conference on Automated Deduction (CADE)*. Springer, Cham, Switzerland, 625–635.

Arthur P. Dempster. 1967. Upper and Lower Probabilities Induced by a Multivalued Mapping. *Annals of Mathematical Statistics* 38, 2 (1967), 325–339.

Dieter Denneberg. 1994. *Non-Additive Measure and Integral*. Kluwer, Dordrecht, Netherlands.

Martin H. Escardó. 2004. Synthetic Topology of Data Types and Classical Spaces. *Electronic Notes in Theoretical Computer Science* 87 (2004), 21–156.

Ronald Fagin and Joseph Y. Halpern. 1991. A New Approach to Updating Beliefs. *Uncertainty in Artificial Intelligence* 6 (1991), 347–374.

Florian Faissole and Bas Spitters. 2017. Synthetic Topology in Homotopy Type Theory for Probabilistic Programming. In *Programming Languages meets Program Verification (PPS)*. Full version with M. E. Bidlingmaier: Faissole–Bidlingmaier–Spitters, “Synthetic topology in Homotopy Type Theory for probabilistic programming,” *Math. Struct. Comput. Sci.* 31(4), 2021..

Tobias Fritz. 2020. A Synthetic Approach to Markov Kernels, Conditional Independence and Theorems on Sufficient Statistics. *Advances in Mathematics* 370 (2020), 107239.

Paolo Ghirardato. 1997. On Independence for Non-Additive Measures, with a Fubini Theorem. *Journal of Economic Theory* 73, 2 (1997), 261–291.

Jean-Yves Girard. 1987. Linear Logic. *Theoretical Computer Science* 50, 1 (1987), 1–101.

Michèle Giry. 1982. A Categorical Approach to Probability Theory. In *Categorical Aspects of Topology and Analysis*. Lecture Notes in Mathematics, Vol. 915. Springer, Berlin, Germany, 68–85.

Andrew M. Gleason. 1957. Measures on the Closed Subspaces of a Hilbert Space. *Journal of Mathematics and Mechanics* 6, 6 (1957), 885–893.

Kurt Gödel. 1933. Eine Interpretation des intuitionistischen Aussagenkalküls. *Ergebnisse eines mathematischen Kolloquiums* 4 (1933), 39–40.

Joseph Y. Halpern. 1999. A Counterexample to Theorems of Cox and Fine. *Journal of Artificial Intelligence Research* 10 (1999), 67–85.

Carl G. Hempel. 1945. Studies in the Logic of Confirmation (I.). *Mind* 54, 213 (1945), 1–26.

Otto Hölder. 1901. Die Axiome der Quantität und die Lehre vom Mass. *Berichte über die Verhandlungen der Sächsischen Gesellschaft der Wissenschaften zu Leipzig* 53 (1901), 1–64.

Kevin S. Van Horn. 2003. Constructing a Logic of Plausible Inference: A Guide to Cox’s Theorem. *International Journal of Approximate Reasoning* 34, 1 (2003), 3–24.

Edwin T. Jaynes. 2003. *Probability Theory: The Logic of Science*. Cambridge University Press, Cambridge, UK.

Xiaodong Jia, Bert Lindenhovius, Michael Mislove, and Vladimir Zamdzhiev. 2021. Commutative Monads for Probabilistic Programming Languages. In *Logic in Computer Science (LICS)*. IEEE.

Peter T. Johnstone. 1982. *Stone Spaces*. Cambridge University Press, Cambridge, UK.

Anders Kock. 2012. Commutative Monads as a Theory of Distributions. *Theory and Applications of Categories* 26 (2012), 97–131.

Andrey N. Kolmogorov. 1933. *Grundbegriffe der Wahrscheinlichkeitsrechnung*. Springer, Berlin, Germany.

The mathlib Community. 2020. The Lean Mathematical Library. In *Certified Programs and Proofs (CPP)*. ACM, New York, NY, USA, 367–381.

J. C. C. McKinsey and Alfred Tarski. 1944. The Algebra of Topology. *Annals of Mathematics* 45, 1 (1944), 141–191.

Jeff B. Paris. 1994. *The Uncertain Reasoner’s Companion: A Mathematical Perspective*. Cambridge University Press, Cambridge, UK.

K. P. S. Bhaskara Rao and M. Bhaskara Rao. 1983. *Theory of Charges: A Study of Finitely Additive Measures*. Pure and Applied Mathematics, Vol. 109. Academic Press, New York, NY, USA.

Giuseppe Rosolini. 1986. *Continuity and Effectiveness in Topoi*. Ph.D. Dissertation. University of Oxford.

Enrique H. Ruspini. 1986. *The Logical Foundations of Evidential Reasoning*. Technical Report 408. SRI International, Menlo Park, CA, USA.

Glenn Shafer. 1976. *A Mathematical Theory of Evidence*. Princeton University Press, Princeton, NJ, USA.

Steven Vickers. 2008. A Localic Theory of Lower and Upper Integrals. *Mathematical Logic Quarterly* 54, 1 (2008), 109–123.

Peter Walley. 1991. *Statistical Reasoning with Imprecise Probabilities*. Chapman and Hall, London, UK.

Brian Weatherston. 2003. From Classical to Intuitionistic Probability. *Notre Dame Journal of Formal Logic* 44, 2 (2003), 111–123.

David H. Wolpert and David B. Kinney. 2024. A Stochastic Model of Mathematics and Science. *Foundations of Physics* 54 (2024), 21.

A Values in the Extended Reals

Valuations take values in the extended non-negative reals $[0, \infty]$, following the measure-theoretic convention of mathlib. As a result, sums indexed by an arbitrary family are unconditionally defined, with no summability hypothesis, and they commute over dependent sums: for a function f on a sigma type $\Sigma a, \beta(a)$, one has $\sum_p f(p) = \sum_a \sum_b f(a, b)$ with no side condition. This is what lets the countable-mixture results of Section 8 hold as stated. Additionally, subtraction is truncated at 0. This never affects the main text, since every subtraction there is applied to a pair already ordered by monotonicity, as with the mass of Section 4, so the truncation is never triggered.

B Lean Statements

B.1 Introduction

B.2 Valuations on Frames

```
structure Valuation (Ω : Type*) [Order.
  Frame Ω] where
  toFun : Ω → ℝ≥0∞
  map_bot' : toFun ⊥ = 0
  map_top' : toFun ⊤ = 1
  mono' : Monotone toFun
  modular' : ∀ a b, toFun a + toFun b =
    toFun (a ⊔ b) + toFun (a ⊓ b)
```

```
theorem Valuation.add_compl_le_one (v :
  Valuation Ω) (a : Ω) :
  v a + v ac ≤ 1
```

```
theorem Valuation.add_compl_eq_sup (v :
  Valuation Ω) (a : Ω) :
  v a + v ac = v (a ⊔ ac)
```

```
theorem Valuation.le_one (v : Valuation Ω)
  (a : Ω) : v a ≤ 1
```

```
theorem Valuation.
  slack_eq_dnGap_add_deMorganGap (v :
  Valuation Ω) (a : Ω) :
  v.slack a = v.dnGap a + v.deMorganGap
  a
```

```
theorem Valuation.slack_eq_sub_add_sub (v
  : Valuation Ω) (a : Ω) :
  v.slack a = (v acc - v a) + (1 - (v a
cc + v ac))
```

```
theorem Valuation.le_compl_compl (v :
  Valuation Ω) (a : Ω) : v a ≤ v acc
```

```
theorem Valuation.slack_submodular (v :
  Valuation Ω) (a b : Ω) :
  v.slack (a ⊔ b) + v.slack (a ⊓ b) ≤ v.
  slack a + v.slack b
```

```
def Order.Frame.IsDeMorgan (Ω : Type*) [
  Order.Frame Ω] : Prop :=
  ∀ a : Ω, acc ⊔ ac = ⊤
```

```
theorem Valuation.
  slack_eq_dnGap_of_isDeMorgan
  (hΩ : Order.Frame.IsDeMorgan Ω) (v :
  Valuation Ω) (a : Ω) :
  v.slack a = v.dnGap a
```

```
noncomputable def Valuation.condVal (v :
  Valuation Ω) (b : Ω) (hb : v b ≠ 0) :
  Valuation Ω
```

```
theorem Valuation.condVal_mul (v :
  Valuation Ω) (b : Ω) (hb : v b ≠ 0) (a
  : Ω) :
  v.condVal b hb a * v b = v (a ⊓ b)
```

```
theorem Valuation.condVal_symm (v :
  Valuation Ω) {a b : Ω}
  (ha : v a ≠ 0) (hb : v b ≠ 0) :
  v.condVal b hb a * v b = v.condVal a
  ha b * v a
```

```
theorem Valuation.cond_add_compl_le (v :
  Valuation Ω) (a b : Ω) :
  v (a ⊓ b) + v (ac ⊓ b) ≤ v b
```

```
theorem Valuation.
  add_compl_eq_one_of_complemented (v :
  Valuation Ω) {a : Ω}
  (ha : a ⊔ ac = ⊤) : v a + v ac = 1
```

```
variable [CompleteBooleanAlgebra Ω]
```

```
theorem classical_additivity (v :
  Valuation Ω) (a : Ω) :
  v a + v ac = 1
```

```
variable [CompleteBooleanAlgebra Ω]
```

```
theorem classical_slack_zero (v :
  Valuation Ω) (a : Ω) :
  v.slack a = 0
```

```
noncomputable def Valuation.mix (w : ι → ℝ
  ≥0∞) (hw : ∑ i, w i = 1)
  (v : ι → Valuation Ω) : Valuation Ω
```

B.3 Where Classicality Lives

```

def Valuation.HasClassicalNegation (v :
  Valuation  $\Omega$ ) : Prop :=
   $\forall a, v\ a + v\ a^c = 1$ 

theorem hasClassicalNegation_of_em (hem :
   $\forall a : \Omega, a \sqcup a^c = \top$ )
  (v : Valuation  $\Omega$ ) : v.
    HasClassicalNegation

theorem em_of_forall_hasClassicalNegation
  (h :  $\forall v : \text{Valuation } \Omega, v.$ 
    HasClassicalNegation) :  $\forall a : \Omega, a$ 
     $\sqcup a^c = \top$ 

theorem sharp_iff_point (v : Valuation  $\Omega$ )
  :
  v.IsSharp  $\leftrightarrow \exists (J : \text{Order.Ideal } \Omega) (hJ$ 
    : J.IsPrime) (htop :  $\top \notin J$ ),
    v = Ideal.toValuation J hJ htop

theorem
  sharp_scott_iff_completelyPrimePoint (v
    : Valuation  $\Omega$ ) :
  v.IsSharp  $\wedge v.$ ScottContinuous  $\leftrightarrow$ 
     $\exists (J : \text{Order.Ideal } \Omega) (hJ : \text{Ideal.}$ 
      IsCompletelyPrime J) (htop :  $\top \notin$ 
      J),
      v = Ideal.toValuation J hJ.1 htop

```

B.4 Valuation Representation

```

noncomputable def Valuation.mass (v :
  Valuation (LowerSet P)) (p : P) :  $\mathbb{R}_{\geq 0\infty}$ 

theorem Valuation.eq_sum_mass (v :
  Valuation (LowerSet P)) (U : LowerSet P
  ) :
  v U =  $\sum p \in U.\text{toFinset}, v.\text{mass } p$ 

theorem Valuation.sum_mass (v : Valuation
  (LowerSet P)) :  $\sum p, v.\text{mass } p = 1$ 

noncomputable def Valuation.toPMF (v :
  Valuation (LowerSet P)) : PMF P

theorem Valuation.exists_pmf_presentation
  { $\Theta : \text{Type } u$ } [Order.Frame  $\Theta$ ] [Finite  $\Theta$ ]
  (v : Valuation  $\Theta$ ) :
   $\exists (P : \text{Type } u) (\_ : \text{PartialOrder } P) (\_$ 
    : Fintype P) ( $\_ : \text{DecidableEq } P$ )
    (e :  $\Theta \simeq \text{LowerSet } P$ ),

```

```

  ( $\forall a, v\ a = \sum p \in (e\ a).\text{toFinset}, (v$ 
    .mapEquiv e).mass p)  $\wedge$ 
     $\sum p, (v.\text{mapEquiv } e).\text{mass } p = 1$ 

noncomputable def deltaPoint (p : P) :
  Valuation (LowerSet P)

theorem Valuation.eq_mix_deltaPoint (v :
  Valuation (LowerSet P)) :
  v = Valuation.mix v.mass v.sum_mass (
    fun p => deltaPoint p)

theorem
  exists_valuation_not_point_representable
  :
   $\exists v : \text{Valuation (LowerSet } \mathbb{N}), \sum' p, v.$ 
    mass p  $\neq v\ \top$ 

theorem topIndicator_mass (n :  $\mathbb{N}$ ) :
  topIndicator.mass n = 0

theorem topIndicator_not_scott :
  ( $\bigsqcup n, \text{topIndicator (LowerSet.Iic } n)$ )  $\neq$ 
    topIndicator  $\top$ 

theorem eq_tsum_mass_of_scott (v :
  Valuation (LowerSet  $\mathbb{N}$ ))
  (hsc :  $v\ \top = \bigsqcup n, v\ (\text{LowerSet.Iic } n)$ )
  :  $v\ \top = \sum' n, v.\text{mass } n$ 

variable [ $\forall p : P, \text{Finite (Set.Iic } p)$ ]

theorem Valuation.isPurelyAtomic_of_scott
  (v : Valuation (LowerSet P))
  (hsc :  $v\ \top = \bigsqcup (S : \text{Finset } P), v\ (S.$ 
    sup LowerSet.Iic)) : IsPurelyAtomic
    v

theorem tsum_mass_le (v : Valuation (
  LowerSet P)) :  $\sum' p, v.\text{mass } p \leq v\ \top$ 

```

B.5 The Dempster–Shafer Reading

```

theorem Valuation.inclusion_exclusion (v :
  Valuation  $\Omega$ ) (s : Finset  $\iota$ )
  (a :  $\iota \rightarrow \Omega$ ) : v (s.sup a) + v.iesEven
    s a = v.iesOdd s a

theorem Valuation.infty_monotone (v :
  Valuation  $\Omega$ ) (s : Finset  $\iota$ ) (a :  $\iota \rightarrow \Omega$ )
  :
  v.iesOdd s a  $\leq v\ ((s.\text{sup } a)^{cc}) + v.$ 
    iesEven s a

```

1981	theorem Valuation.two_monotone (v :	($\forall x y : \Omega, x \sqcap y = \perp \rightarrow q(x \sqcup y) =$	2036
1982	Valuation Ω) (a b : Ω) :	$q x + q y) \wedge$	2037
1983	v a + v b $\leq$ v (a $\sqcup$ b) ^{cc} + v (a $\sqcap$ b)	$\neg (\forall x y : \Omega, q x + q y = q(x \sqcup y)$	2038
1984		+ q (x $\sqcap$ y))	2039
1985	theorem Valuation.plausibility_sub_self (v		2040
1986	: Valuation Ω) (a : Ω) :	theorem no_disjunction_functional :	2041
1987	v.plausibility a - v a = v.slack a	$\neg \exists S : \mathbb{R}_{\geq 0}^{\infty} \rightarrow \mathbb{R}_{\geq 0}^{\infty} \rightarrow \mathbb{R}_{\geq 0}^{\infty},$	2042
1988		$\forall x y : \text{LowerSet } V, qV(x \sqcup y) = S$	2043
1989	theorem Valuation.sup_compl_decomp (v :	(qV x) (qV y)	2044
1990	Valuation Ω) (a b : Ω) :		2045
1991	v (a $\sqcup$ b ^c) = v b ^c + (v (a $\sqcap$ b) + v.	theorem constructive_cox (M : CoxModel Ω)	2046
1992	emGap a b)	(hmod : $\forall a b : \Omega, M.\text{pl } a \top + M.\text{pl } b \top$	2047
1993		= M.pl (a $\sqcup$ b) $\top$ + M.pl (a $\sqcap$ b) $\top$)	2048
1994	noncomputable def Valuation.dempsterCond (	:	2049
1995	v : Valuation Ω) (b : Ω)	$\exists (v : \text{Valuation } \Omega) (g : \mathbb{R} \rightarrow \mathbb{R}_{\geq 0}^{\infty}),$	2050
1996	(hb : v b $\neq$ 0) : Valuation Ω	StrictMonoOn g (Set.Icc 0 1) $\wedge$ g 0 =	2051
1997		0 $\wedge$ g 1 = 1 $\wedge$	2052
1998	theorem Valuation.dempsterCond_apply (v :	$\forall a, v a = g(M.\text{pl } a \top)$	2053
1999	Valuation Ω) (b : Ω) (hb : v b $\neq$ 0)		2054
2000	(a : Ω) : v.dempsterCond b hb a = (v (	noncomputable def coxModelENNReal :	2055
2001	a $\sqcup$ b ^c) - v b ^c) / (1 - v b ^c)	CoxModel $\mathbb{R}_{\geq 0}^{\infty}$	2056
2002			2057
2003	theorem Valuation.	theorem Scale.exists_mul_generator :	2058
2004	dempsterCond_eq_condVal_iff_slack (v :	$\exists G : \mathbb{R} \rightarrow \mathbb{R},$	2059
2005	Valuation Ω) {b : Ω }	($\forall x y, S.u \leq x \rightarrow S.u \leq y \rightarrow G(S.F x$	2060
2006	(hb : v b $\neq$ 0) :	y) = G x * G y) $\wedge$	2061
2007	($\forall a, v.\text{dempsterCond } b \text{ hb } a = v.$	($\forall x y, S.u \leq x \rightarrow x < y \rightarrow G x < G y$)	2062
2008	condVal b hb a) $\leftrightarrow$ v.slack b = 0		2063
2009		theorem Scale.aczelStatement_cone :	2064
2010		$\exists G : \mathbb{R} \rightarrow \mathbb{R},$	2065
2011	B.6 Cox's Theorem	($\forall x y, S.u \leq x \rightarrow S.u \leq y \rightarrow G(S.F x$	2066
2012		y) = G x * G y) $\wedge$	2067
2013	structure CoxModel ($\Omega : \text{Type}^*$) [Order.	StrictMonoOn G (Set.Ici S.u)	2068
2014	Frame Ω] where		2069
2015	pl : $\Omega \rightarrow \Omega \rightarrow \mathbb{R}$	theorem exists_diag_eq {F : $\mathbb{R} \rightarrow \mathbb{R} \rightarrow \mathbb{R}$ } {a	2070
2016	mono_left : $\forall c, \text{Monotone fun } a \Rightarrow \text{pl } a$	b c : $\mathbb{R}$ } (hab : a $\leq$ b)	2071
2017	c	(hcont : ContinuousOn (fun x => F x x)	2072
2018	pl_bot : $\forall c, \text{pl } \perp c = 0$	(Set.Icc a b))	2073
2019	pl_top : $\forall c, \text{pl } \top c = 1$	(hc : c $\in$ Set.Icc (F a a) (F b b)) : $\exists$	2074
2020	F : $\mathbb{R} \rightarrow \mathbb{R} \rightarrow \mathbb{R}$	x $\in$ Set.Icc a b, F x x = c	2075
2021	conj : $\forall a b c, \text{pl } (a \sqcap b) c = F(\text{pl } a ($		2076
2022	b $\sqcap$ c)) (pl b c)	structure Scale where	2077
2023	F_assoc : $\forall x y z, F(F x y) z = F x (F$	F : $\mathbb{R} \rightarrow \mathbb{R} \rightarrow \mathbb{R}$	2078
2024	y z)	hdiv : $\forall c, \exists x, F x x = c$	2079
2025	F_strictMono_left : $\forall y, 0 < y \rightarrow$	hassoc : $\forall x y z, F(F x y) z = F x (F y$	2080
2026	StrictMono fun x => F x y	z)	2081
2027		hpos : $\forall c x, x < F x c$	2082
2028	theorem isDeMorgan_lowerSetV : Order.Frame	hmono : $\forall x_1 x_2 y_1 y_2 : \mathbb{R}, x_1 \leq x_2 \rightarrow y_1$	2083
2029	.IsDeMorgan (LowerSet V)	$\leq y_2 \rightarrow F x_1 y_1 \leq F x_2 y_2$	2084
2030		hcont : Continuous fun p : $\mathbb{R} \times \mathbb{R} \Rightarrow F p$	2085
2031	theorem modularity_irreducible :	.1 p.2	2086
2032	$\exists (\Omega : \text{Type}) (_ : \text{Order.Frame } \Omega) (q :$	hbot : $\forall x, \text{Filter.Tendsto (fun } z \Rightarrow F x$	2087
2033	$\Omega \rightarrow \mathbb{R}_{\geq 0}^{\infty}),$	z) Filter.atBot (nhds x)	2088
2034	Monotone q $\wedge$ q $\perp = 0 \wedge$ q $\top = 1 \wedge$	u : $\mathbb{R}$	2089
2035			2090

B.7 Two Grounding Models

```

noncomputable def MeasureTheory.Measure.
  toValuationOpens
  (μ : Measure X) [IsProbabilityMeasure
    μ] : Valuation (Opens X)

theorem toValuationOpens_slack_eq_frontier
  (μ : Measure X)
  [IsProbabilityMeasure μ] (U : Opens X)
  :
  μ.toValuationOpens.slack U = μ (
    frontier (U : Set X))

theorem classical_posterior_compl_eq (μ :
  Measure X)
  [IsProbabilityMeasure μ] (a b : Opens
    X)
  (hb : μ.toValuationOpens b ≠ 0) :
  μ ((a : Set X)c â⊗ b) / μ b
  = μ.toValuationOpens.condVal b hb ac
  + μ (frontier (a : Set X) â⊗ b) /
    μ b

theorem interiorMeasure_add_compl_le (μ :
  Measure X) [IsProbabilityMeasure μ]
  (S : Set X) : μ.interiorMeasure S + μ.
    interiorMeasure Sc ≤ 1

theorem
  isDeMorgan_opens_iff_extremallyDisconnected
  :
  Order.Frame.IsDeMorgan (Opens X) ↔
    ExtremallyDisconnected X

theorem isDeMorgan_opens_iff_stonean [
  CompactSpace X] [T2Space X] :
  Order.Frame.IsDeMorgan (Opens X) ↔
    ExtremallyDisconnected X

inductive ThreePt : Type
  | ptA : ThreePt
  | ptB : ThreePt
  | ptZ : ThreePt

theorem frontier_subset_ptZ {U : Set
  ThreePt} (hU : IsOpen U) :
  frontier U ⊆ {ptZ}

theorem exists_ne_valuation_same_slack :
  ∃ v v' : Valuation (Opens ThreePt), v.
    slack = v'.slack ∧ v ≠ v'

```

```

theorem haltsOpen_compl : haltsOpenc = (⊥
  : Opens Prop)

theorem slack_eq_boundary {p : ℝ≥0∞} (hp1
  : p ≤ 1) :
  (sierpinskiValuation hp1).slack
    haltsOpen
  = coinMeasure p (frontier {q : Prop
    | q})

theorem sharpReadout_not_computable (n : ℕ
  ) :
  ¬ComputablePred fun c : Code =>
    sharpReadout n c haltsOpen = 1

theorem isOpen_wordCylinder (w : List Bool
  ) : IsOpen (wordCylinder w)

theorem cantorMeasure_wordCylinder (w :
  List Bool) :
  cantorMeasure (wordCylinder w) = (1 /
    2) ^ w.length

theorem
  disjoint_wordCylinder_of_not_prefix {v
  w : List Bool}
  (hv : ¬ v <+: w) (hw : ¬ w <+: v) :
  Disjoint (wordCylinder v) (
    wordCylinder w)

theorem cantorMeasure_biUnion_wordCylinder
  {F : Finset (List Bool)}
  (hF : PrefixFree F) :
  cantorMeasure (⋃ w ∈ F, wordCylinder w
  ) = ∑ w ∈ F, (1 / 2 : ℝ≥0∞) ^ w.
    length

```

B.8 The Monad Question

```

def Valuation.comap (f : FrameHom Γ Ω) (v
  : Valuation Ω) : Valuation Γ

theorem Valuation.comap_id (v : Valuation
  Ω) : v.comap (FrameHom.id Ω) = v

theorem Valuation.comap_comp (f : FrameHom
  Γ Ω) (g : FrameHom Δ Γ)
  (v : Valuation Ω) : v.comap (f.comp g)
  = (v.comap f).comap g

theorem exists_rect_decomposition (W :
  LowerSet (ℕ × ℕ)) :
  ∃ (N : ℕ) (A B : ℕ → LowerSet ℕ),

```

```

2201     W = (Finset.range N).sup fun i => A
2202         i ×s B i
2203
2204 theorem Valuation.eq_on_finsetSup_of_eq_on
2205   {G : Set Ω}
2206   (hG : ∀ g ∈ G, ∀ g' ∈ G, g ⊔ g' ∈ G) (
2207     u u' : Valuation Ω)
2208   (h : ∀ g ∈ G, u g = u' g) {ι : Type*}
2209     [DecidableEq ι]
2210   (s : Finset ι) (a : ι → Ω) (ha : ∀ i ∈
2211     s, a i ∈ G) :
2212     u (s.sup a) = u' (s.sup a)
2213
2214 theorem Valuation.product_unique (v w :
2215   Valuation (LowerSet ℕ))
2216   {u u' : Valuation (LowerSet (ℕ × ℕ))}
2217   (hu : ∀ A B : LowerSet ℕ, u (A ×s B) =
2218     v A * w B)
2219   (hu' : ∀ A B : LowerSet ℕ, u' (A ×s B)
2220     = v A * w B) : u = u'
2221
2222 noncomputable def ultrafilterValuation (W
2223   : Ultrafilter ι) : Valuation (Set ι)
2224
2225 theorem rect_mem_iff {α β : Type*} {W :
2226   Ultrafilter (α × β)}
2227   {S : Set α} {T : Set β} :
2228   S ×s T ∈ W ↔ S ∈ W.map Prod.fst ∧ T ∈
2229     W.map Prod.snd
2230
2231 theorem product_valuation_not_unique :
2232   ∃ u u' : Valuation (Set (ℕ × ℕ)),
2233     (∀ S T : Set ℕ, u (S ×s T) = u' (S ×
2234       s T)) ∧ u ≠ u'
2235
2236 theorem fubini_fails_for_valuations :
2237   ∃ (v w : Valuation (Set ℕ)) (u u' :
2238     Valuation (Set (ℕ × ℕ))),
2239     (∀ S T : Set ℕ, u (S ×s T) = v S * w
2240       T) ∧
2241     (∀ S T : Set ℕ, u' (S ×s T) = v S *
2242       w T) ∧ u ≠ u'
2243
2244 noncomputable def Valuation.mixCountable {
2245   ι : Type*} (w : ι → ℝ≥0)
2246   (hw : ∑' i, w i = 1) (v : ι →
2247     Valuation Ω) : Valuation Ω
2248
2249 theorem Valuation.
2250   mixCountable_mixCountable {ι : Type*} {
2251     κ : ι → Type*}
2252   (w : ι → ℝ≥0) (hw : ∑' i, w i = 1)
2253   (u : ∀ i, κ i → ℝ≥0) (hu : ∀ i, ∑' j,
2254     u i j = 1)
2255

```

```

2256   (v : ∀ i, κ i → Valuation Ω) :
2257   Valuation.mixCountable w hw
2258     (fun i => Valuation.mixCountable (
2259       u i) (hu i) (v i))
2260   = Valuation.mixCountable (fun p : Σ
2261     i, κ i => w p.1 * u p.1 p.2)
2262     (tsum_sigma_mul_eq_one w hw u hu
2263       ) (fun p => v p.1 p.2)
2264
2265 theorem Valuation.mixCountable_unique (v :
2266   Valuation Ω) :
2267   Valuation.mixCountable (fun _ : PUnit
2268     => 1) (by simp) (fun _ => v) = v
2269
2270 theorem Valuation.mixCountable_comm {ι κ :
2271   Type*}
2272   (w : ι → ℝ≥0) (hw : ∑' i, w i = 1) (u
2273     : κ → ℝ≥0) (hu : ∑' j, u j = 1)
2274   (v : ι → κ → Valuation Ω) :
2275   Valuation.mixCountable w hw (fun i =>
2276     Valuation.mixCountable u hu (v i))
2277   = Valuation.mixCountable u hu
2278     (fun j => Valuation.mixCountable
2279       w hw (fun i => v i j))
2280

```

B.9 Valuations Beyond Frames

Theorem 52.

```

2284 theorem exists_sharp_separating {a b : Ω}
2285   (hab : ¬ a ≤ b) :
2286   ∃ v : Valuation Ω, v.IsSharp ∧ v a = 1
2287     ∧ v b = 0
2288

```

Definition 50.

```

2290 structure LValuation (Ω : Type*) [Lattice
2291   Ω] [BoundedOrder Ω] where
2292   toFun : Ω → ℝ≥0
2293   map_bot' : toFun ⊥ = 0
2294   map_top' : toFun ⊤ = 1
2295   mono' : Monotone toFun
2296   modular' : ∀ a b, toFun a + toFun b
2297     = toFun (a ⊔ b) + toFun (a ⊓
2298       b)
2299

```

Proposition 51.

```

2301 def LValuation.comap (f :
2302   BoundedLatticeHom Γ Ω) (v : LValuation
2303     Ω) :
2304   LValuation Γ
2305
2306 theorem LValuation.comap_id (v :
2307   LValuation Ω) :
2308   v.comap (BoundedLatticeHom.id Ω) = v
2309

```

```

2311 theorem LValuation.comap_comp (f :
2312   BoundedLatticeHom  $\Gamma$   $\Omega$ )
2313   (g : BoundedLatticeHom  $\Delta$   $\Gamma$ ) (v :
2314     LValuation  $\Omega$ ) :
2315     v.comap (f.comp g) = (v.comap f).comap
2316       g

```

Theorem 53.

```

2319 theorem LValuation.exists_sharp_separating
2320   [DistribLattice  $\Omega$ ]
2321   {a b :  $\Omega$ } (hab :  $\neg a \leq b$ ) :
2322      $\exists v : \text{LValuation } \Omega, v.\text{IsSharp} \wedge v a =$ 
2323       1  $\wedge v b = 0$ 

```

Definition 54.

```

2325 inductive M3 : Type
2326   | z | a | b | c | o
2327
2328 instance : Lattice M3 where
2329   -- z bottom, o top, a b c pairwise-
2330     incomparable atoms;
2331   -- any two distinct atoms join to o and
2332     meet to z
2333
2334 instance : BoundedOrder M3 where
2335   top := o
2336   bot := z
2337
2338 theorem M3.not_distributive :
2339    $\neg \forall x y w : M3, x \sqcap (y \sqcup w) = (x \sqcap y)$ 
2340      $\sqcup (x \sqcap w)$ 

```

Theorem 55.

```

2343 theorem M3.no_sharp_valuation (v :
2344   LValuation M3) :  $\neg v.\text{IsSharp}$ 

```

Definition 56.

```

2347 def IsTensorValuation (v :  $\Omega \rightarrow \mathbb{R}_{\geq 0\infty}$ ) (t :
2348    $\Omega \rightarrow \Omega \rightarrow \Omega$ ) : Prop :=
2349    $\forall a b, v a + v b = v (a \sqcup b) + v (t a b)$ 
2350

```

Theorem 57.

```

2352 theorem IsTensorValuation.collapse {v :  $\Omega$ 
2353    $\rightarrow \mathbb{R}_{\geq 0\infty}$ } {t :  $\Omega \rightarrow \Omega \rightarrow \Omega$ }
2354   (h : IsTensorValuation v t) {a :  $\Omega$ } (
2355     ha :  $v a \neq \top$ ) :
2356     v (t a a) = v a

```

Example 58.

```

2359 theorem forced_constant_of_add_valuation {
2360   v :  $\mathbb{N} \rightarrow \mathbb{R}_{\geq 0\infty}$ }
2361   (h : IsTensorValuation v ( $\hat{\cdot} + \hat{\cdot}$ )) (
2362     hfin :  $\forall n, v n \neq \top$ ) (n :  $\mathbb{N}$ ) :
2363     v n = v 0

```

Corollary 59.

```

theorem
  not_isTensorValuation_of_discriminates
  {v :  $\Omega \rightarrow \mathbb{R}_{\geq 0\infty}$ }
  {t :  $\Omega \rightarrow \Omega \rightarrow \Omega$ } {a :  $\Omega$ } (ha :  $v a \neq \top$ )
  (hne :  $v (t a a) \neq v a$ ) :
   $\neg \text{IsTensorValuation } v t$ 

```

C The Aczél Development

This appendix sketches the from-scratch construction behind the product-rule half of Theorem 30. The target shape is Aczél's associativity theorem [Aczél 1966]: a suitably regular associative $F : \mathbb{R} \rightarrow \mathbb{R} \rightarrow \mathbb{R}$ reg graduates to multiplication, $G(F x y) = G(x) \cdot G(y)$ for all x, y and a strictly monotone G . mathlib has no associativity representation, no Hölder-style embedding of ordered semigroups on intervals, and no Cauchy functional equation under monotonicity (only under continuity).

Hypotheses. A Scale bundles: associativity of F , *divisibility* $\forall c \exists x, F x x = c$, positivity on the cone $x < F x c$, joint order-preservation, joint continuity, a unit u , and identity at $-\infty$, meaning $F x z \rightarrow x$ as $z \rightarrow -\infty$. Order-preservation is required by the additivity proof below. Divisibility yields the 2^n -th roots used in the dyadic construction below. Without it the generator is not definable. It is assumed directly in the bundle, though its interval form follows from the intermediate value theorem (exists_diag_eq). Non-vacuity: *log-sum-exp*, $F x y = \log(e^x + e^y)$ with $u = 0$, satisfies every field of the bundle (logSumExpScale), with intended generator globally exp.

Construction. Write $a^{(n)}$ for n right-associated copies of a under F (Fpow). Associativity alone makes $n \mapsto a^{(n)}$ a semigroup morphism from $(\mathbb{N}, +)$. Commutativity is not assumed, and follows instead as a corollary of the generator. Unboundedness of the iterates is derived: were they bounded, monotone convergence would give a fixed point $F L a = L$, contradicting positivity, and unboundedness yields the Archimedean property. Divisibility yields a tower of 2^n -th roots $u \succ u^{1/2} \succ u^{1/4} \succ \dots$ (droot), where 2^ℓ copies of the finer root compose under F to one copy of the coarser root. Counting the number of copies of the n -th root not exceeding b (dcount, well defined by the Archimedean property) gives dyadic approximants $g_n(b) = \text{dcount}(b, n)/2^n$. The refinement inequality $2N + 1 \leq N' \leq 2N + 2$, where N and N' are the counts at consecutive refinement levels n and $n + 1$, makes (g_n) monotone and Cauchy, and the generator is its limit, $g(b) = \bigsqcup_n g_n(b)$.

Results (all sorry-free, standard axioms). On the cone $\{x \mid u \leq x\}$ we prove three things. Normalization: $g(u) = 1$. Additivity: $g(F x y) = g(x) + g(y)$, since $\text{dcount } x + \text{dcount } y + 1 \leq \text{dcount}(F x y) \leq$

dcount x + dcount y + 3 and the 2^{-n} error term vanishes. And strict monotonicity, proved without any continuity of g : a root inserted between $x < y$ contributes a full extra block of 2^{n-M} copies, M being the refinement level of the inserted root. Exponentiating gives the multiplicative form, a strictly monotone G with $G(F x y) = G(x) G(y)$ on the cone (exists_mul_generator), the conclusion of Aczél's theorem restricted to the cone (aczelStatement_cone). The reorientation $\bar{G} = \exp(-g)$ gives the $[0, 1]$ -valued form used in Section 6.

Remaining gap. The constructed g is discontinuous at the unit (it is 0 below u), so matching the classical statement verbatim needs the off-cone extension (a group completion) and continuity of the extension. This is open. It is not vacuous: the log-sum-exp example admits a global continuous generator (hasOrderedGenerator_logSumExp). Of Hölder's classical hypotheses, the Archimedean property is proved, and is what enables the embedding into $\mathbb{R}$. The interval form of divisibility follows from the intermediate value theorem. Only the embedding's global regularity remains.